

Software-level Attacks on Architectural and Microarchitectural State



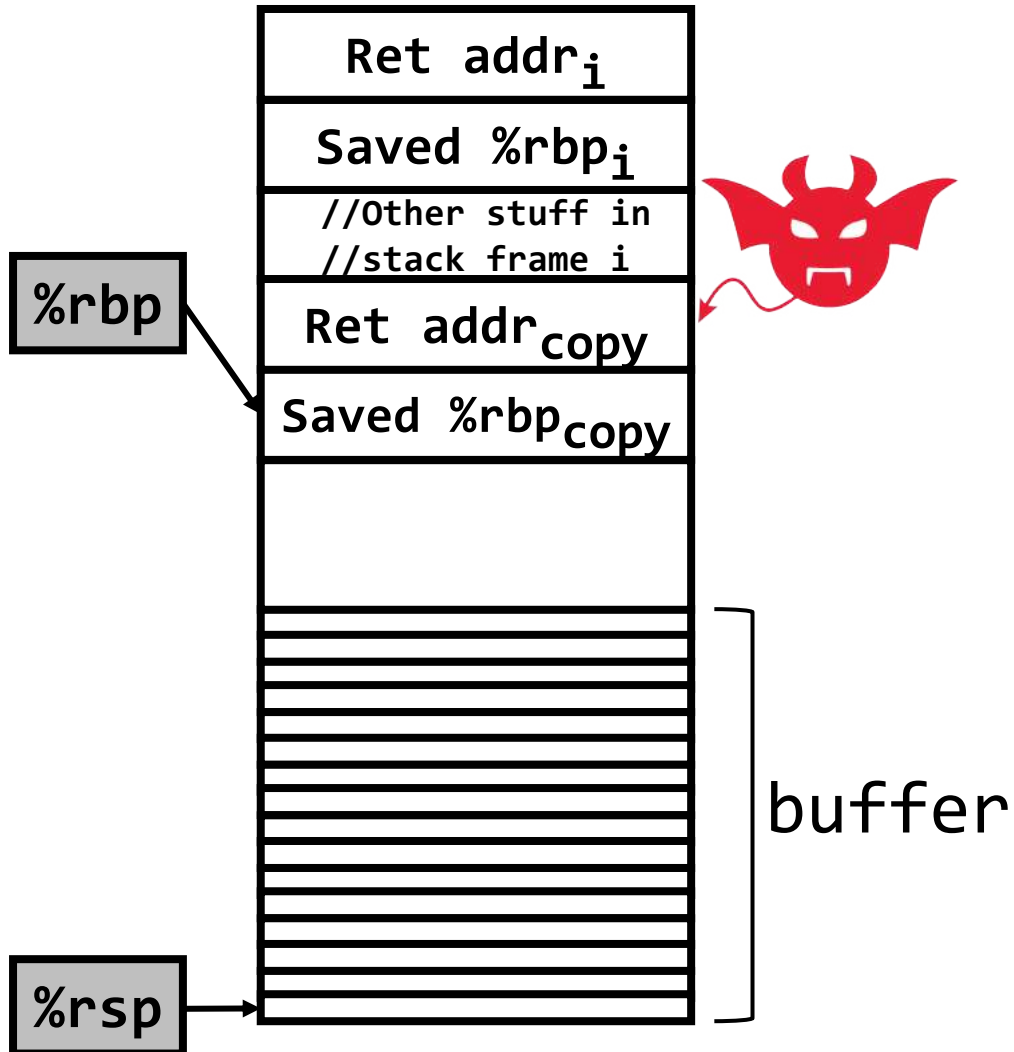
ACACES SUMMER SCHOOL

July 9th - 10th, 2020

~~Fiuggi, Rome~~ Zoom World

James Mickens
Harvard University

Summary of the Last Lecture



- Buffer overflows let attackers overwrite code pointers in the path of the overflow
- Subverting control flow allows:
 - Return to attacker-provided shellcode
 - Return to **libc** (or another preexisting function)

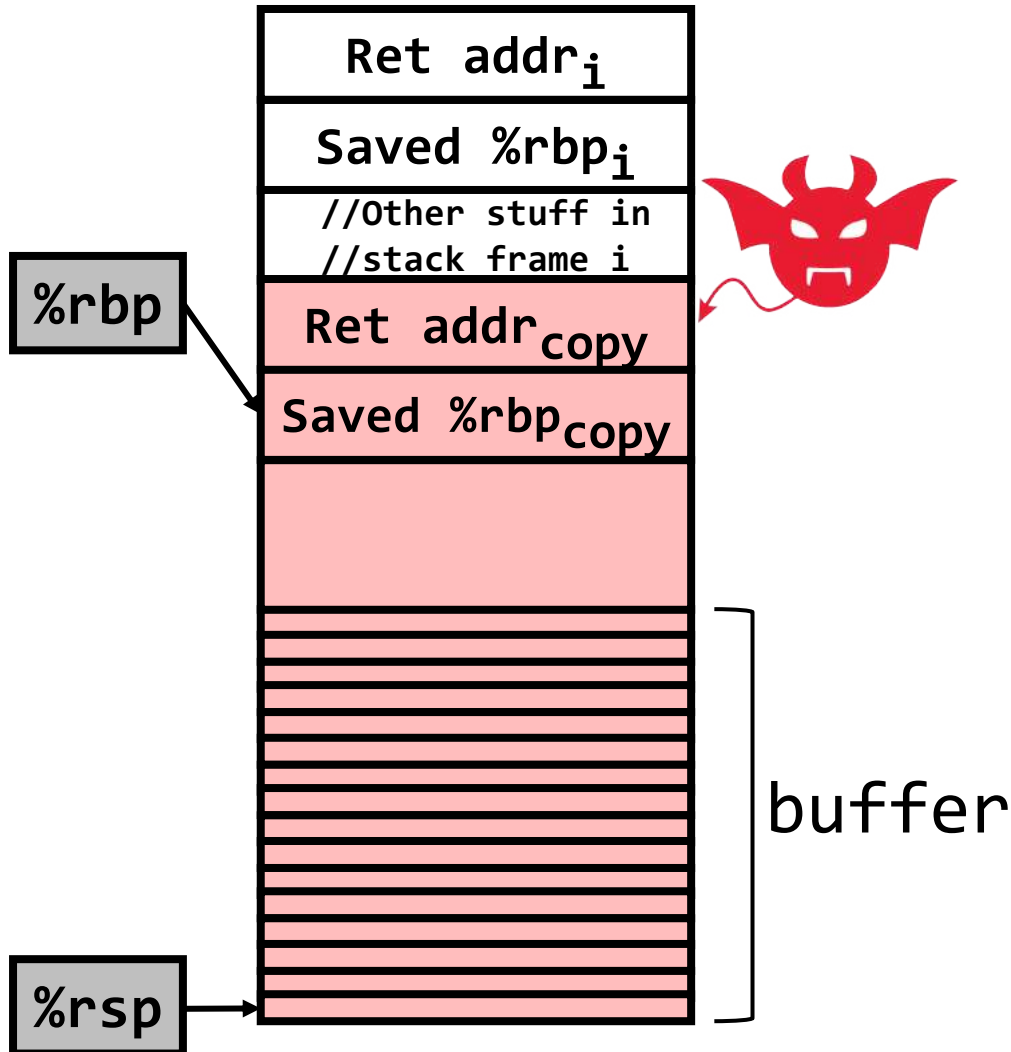
Summary of the Last Lecture

- Defenses:

- Canaries: Detect corruption of return addresses
- NX bits: Prevent injection of shellcode
- ASLR + struct randomization: Prevent attacker from easily locating program state

- Given these defenses, attackers require mechanisms to:

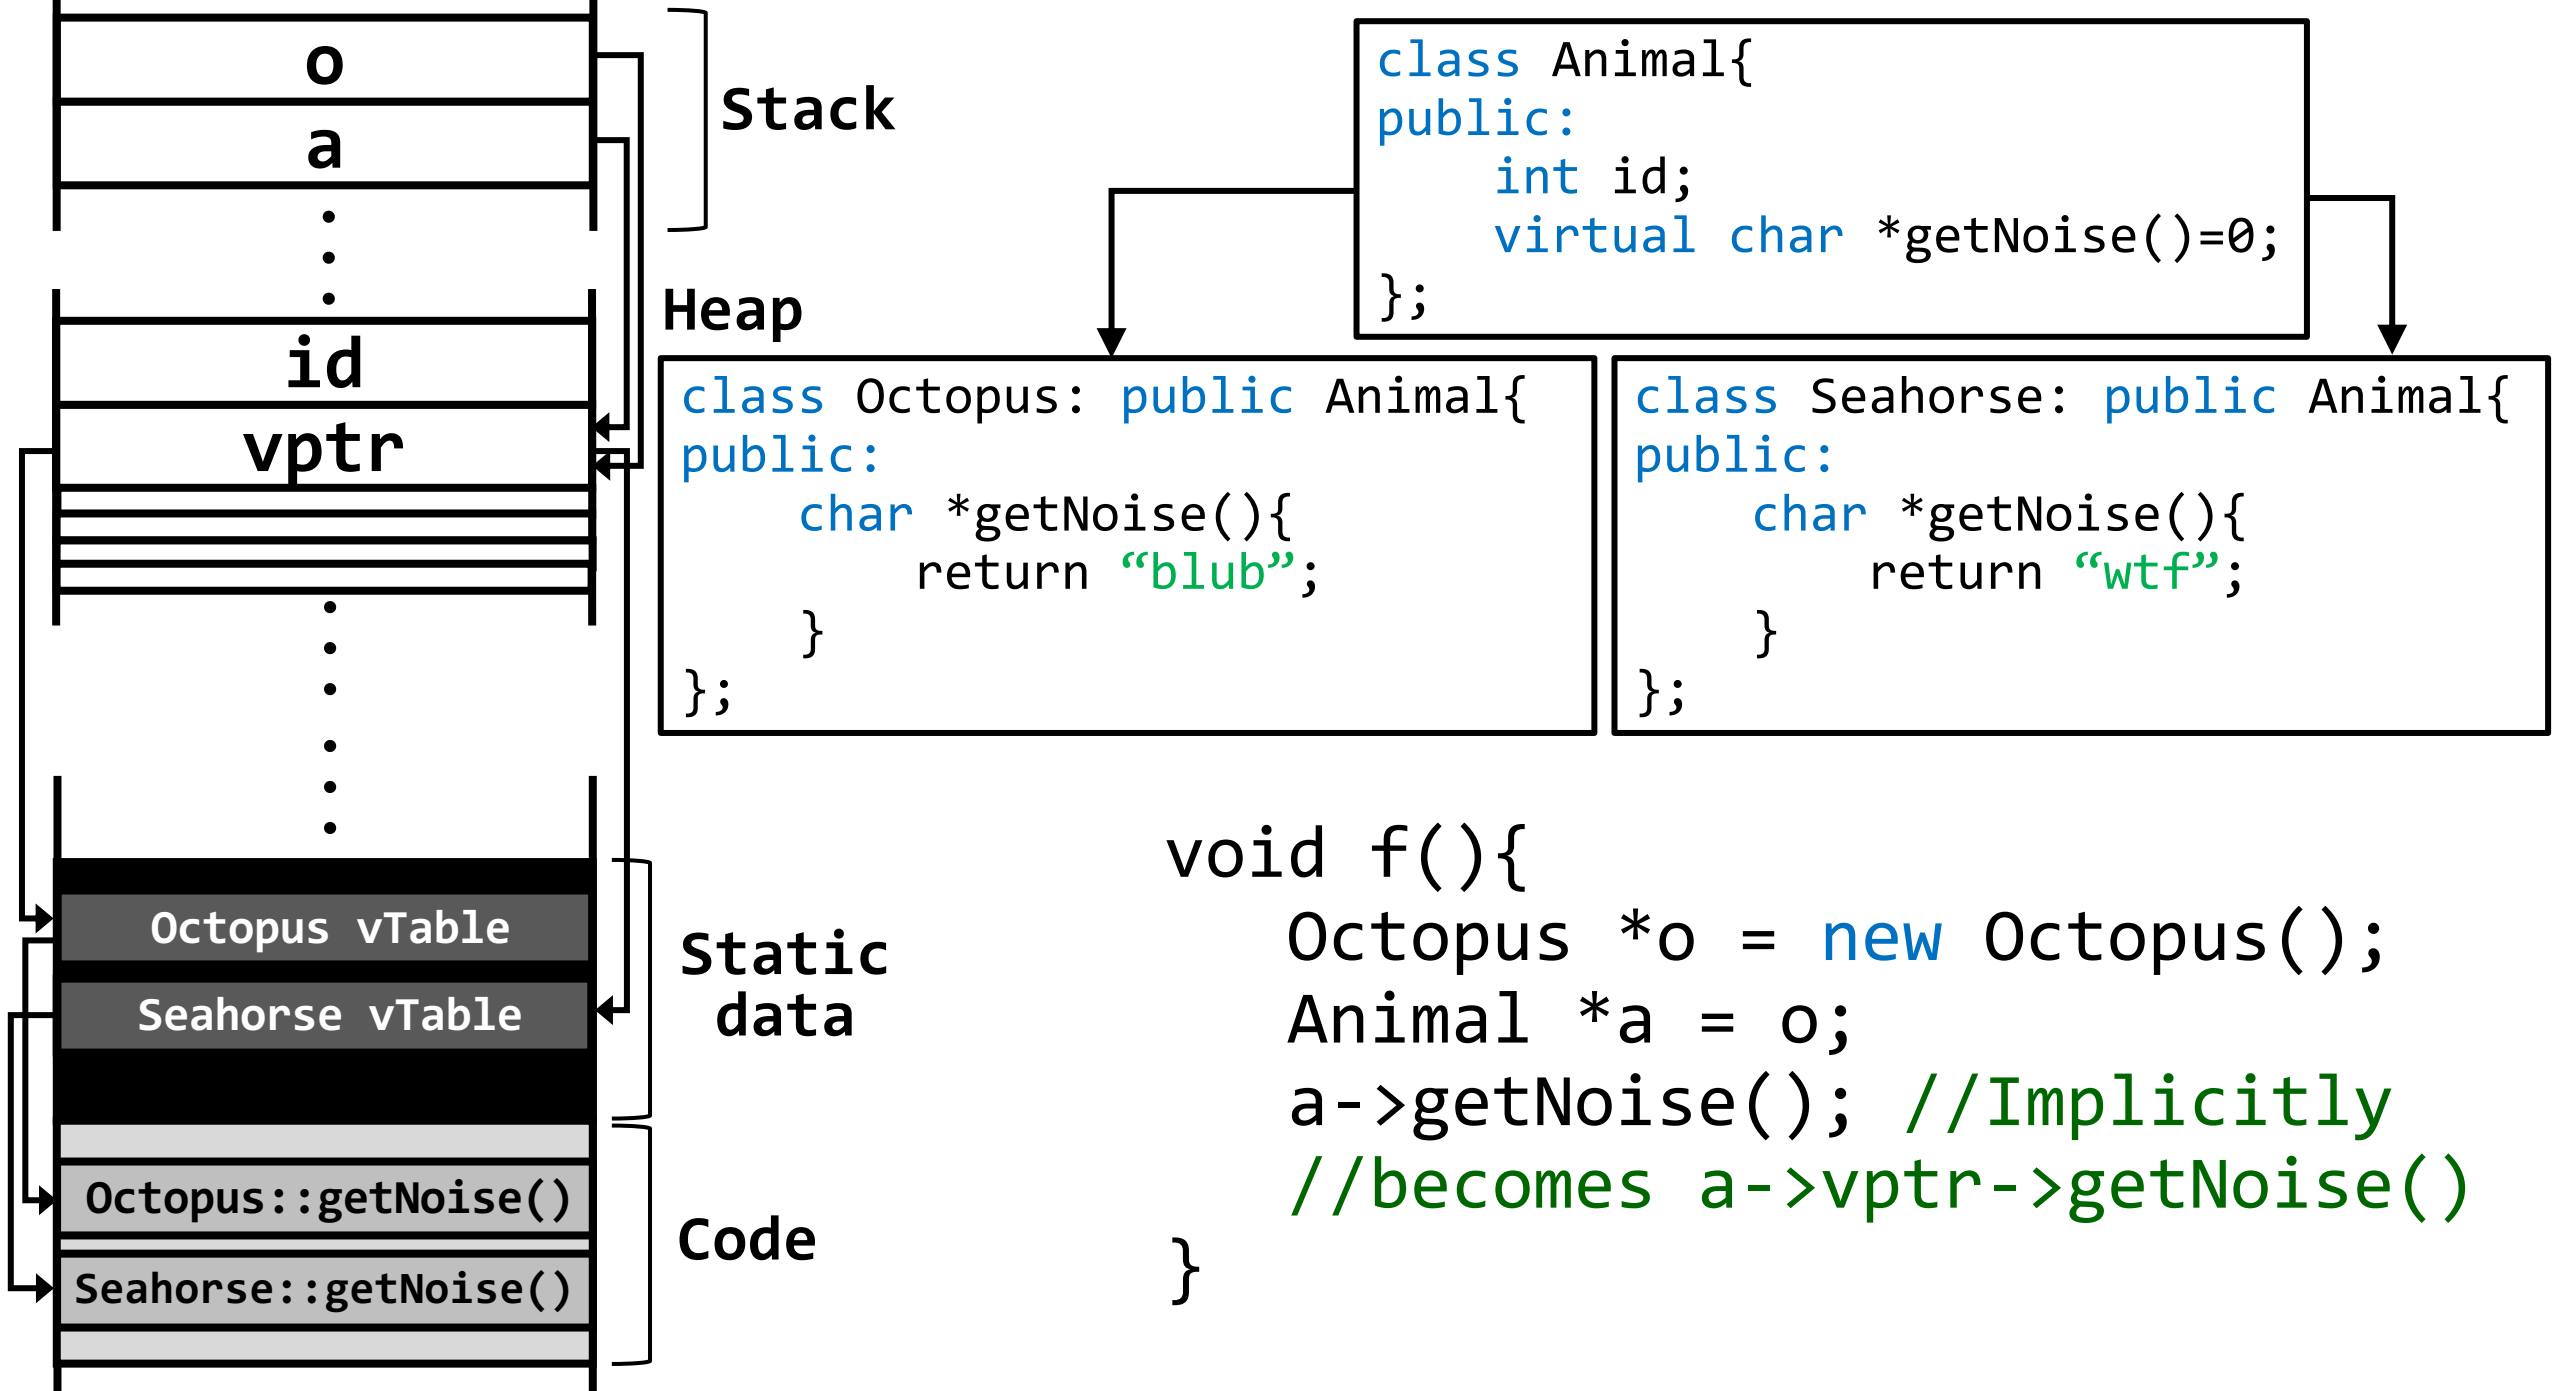
- Read memory (to locate program state)
- Write memory in a non-contiguous fashion (to avoid corrupting canaries)
- Implement arbitrary malicious behavior using only preexisting code (to escape the shackles of NX bits)




The Magic of Vptrs



- Modern languages like C++, Go, Java, C#, and Rust allow the binding of an abstract method declaration to multiple concrete implementations
 - C++ polymorphism
 - Interfaces in Go, Java, and C#
- Bindings are enabled through a layer of indirection
- By corrupting the indirection tables, the attacker can create memory exploits (both read and write)!





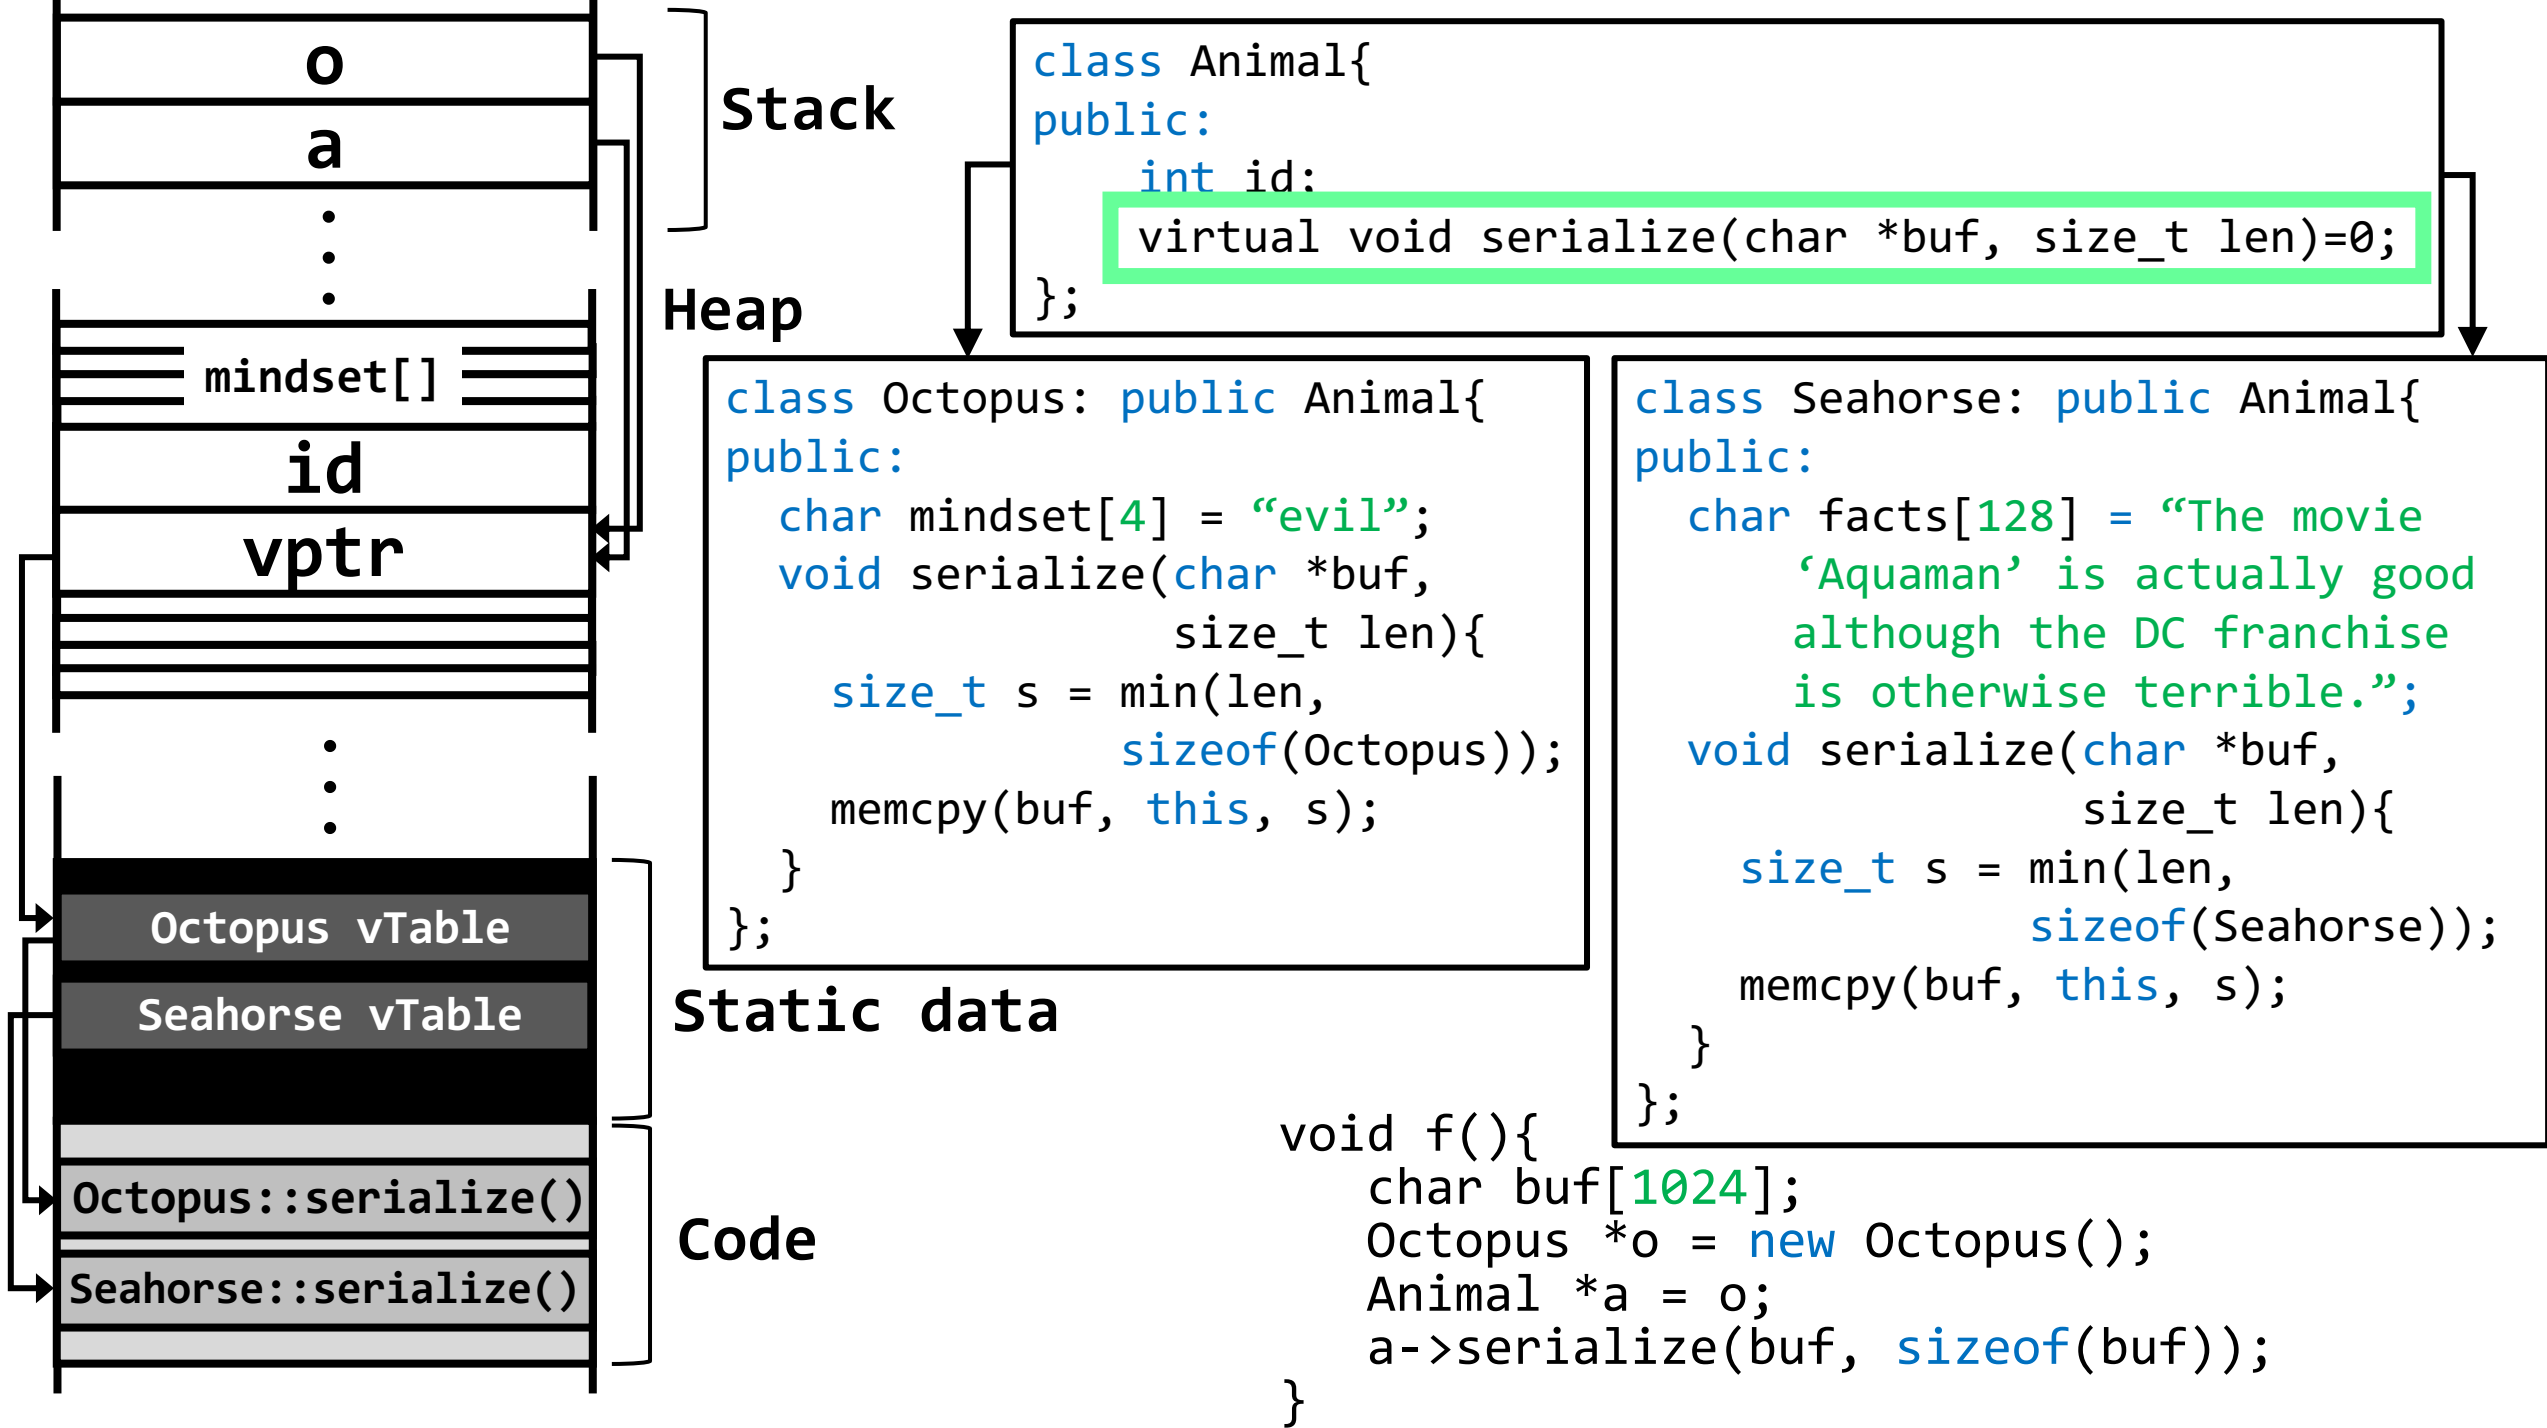
TYPE CONFUSION

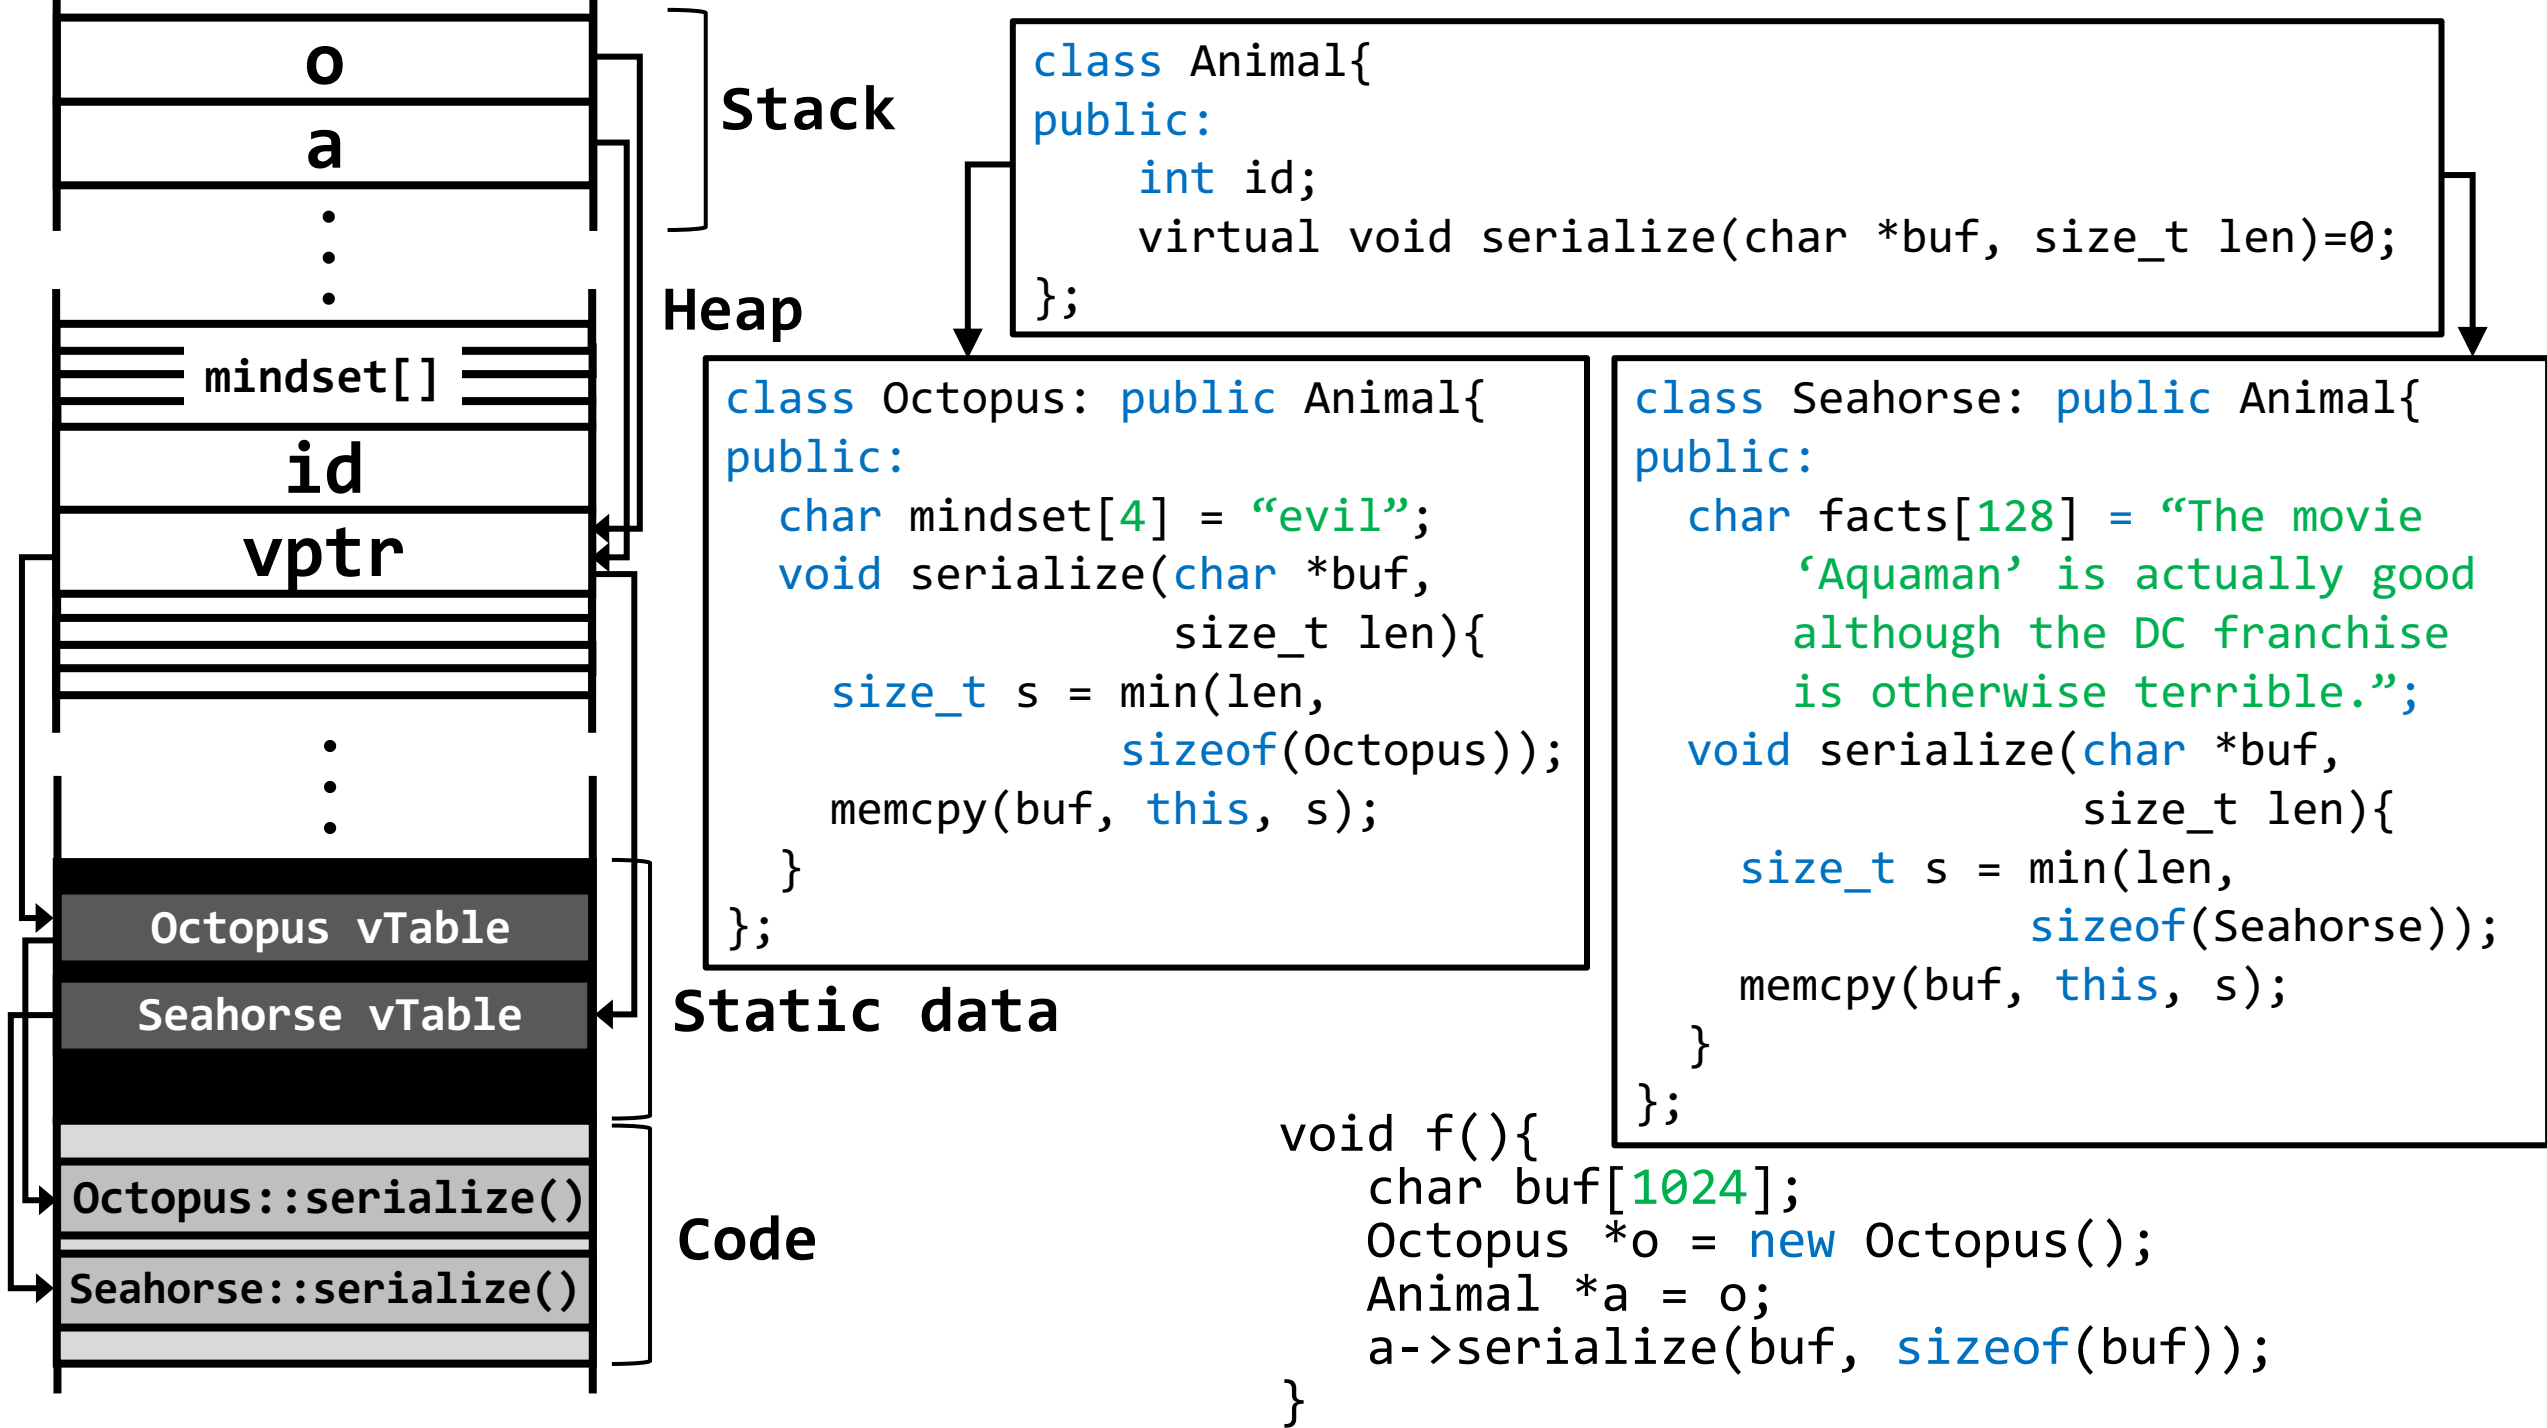
```
void f(){
    Octopus *o = new Octopus();
    Animal *a = o;
    a->getNoise(); //Implicitly
                  //becomes a->vptr->getNoise()
}
```

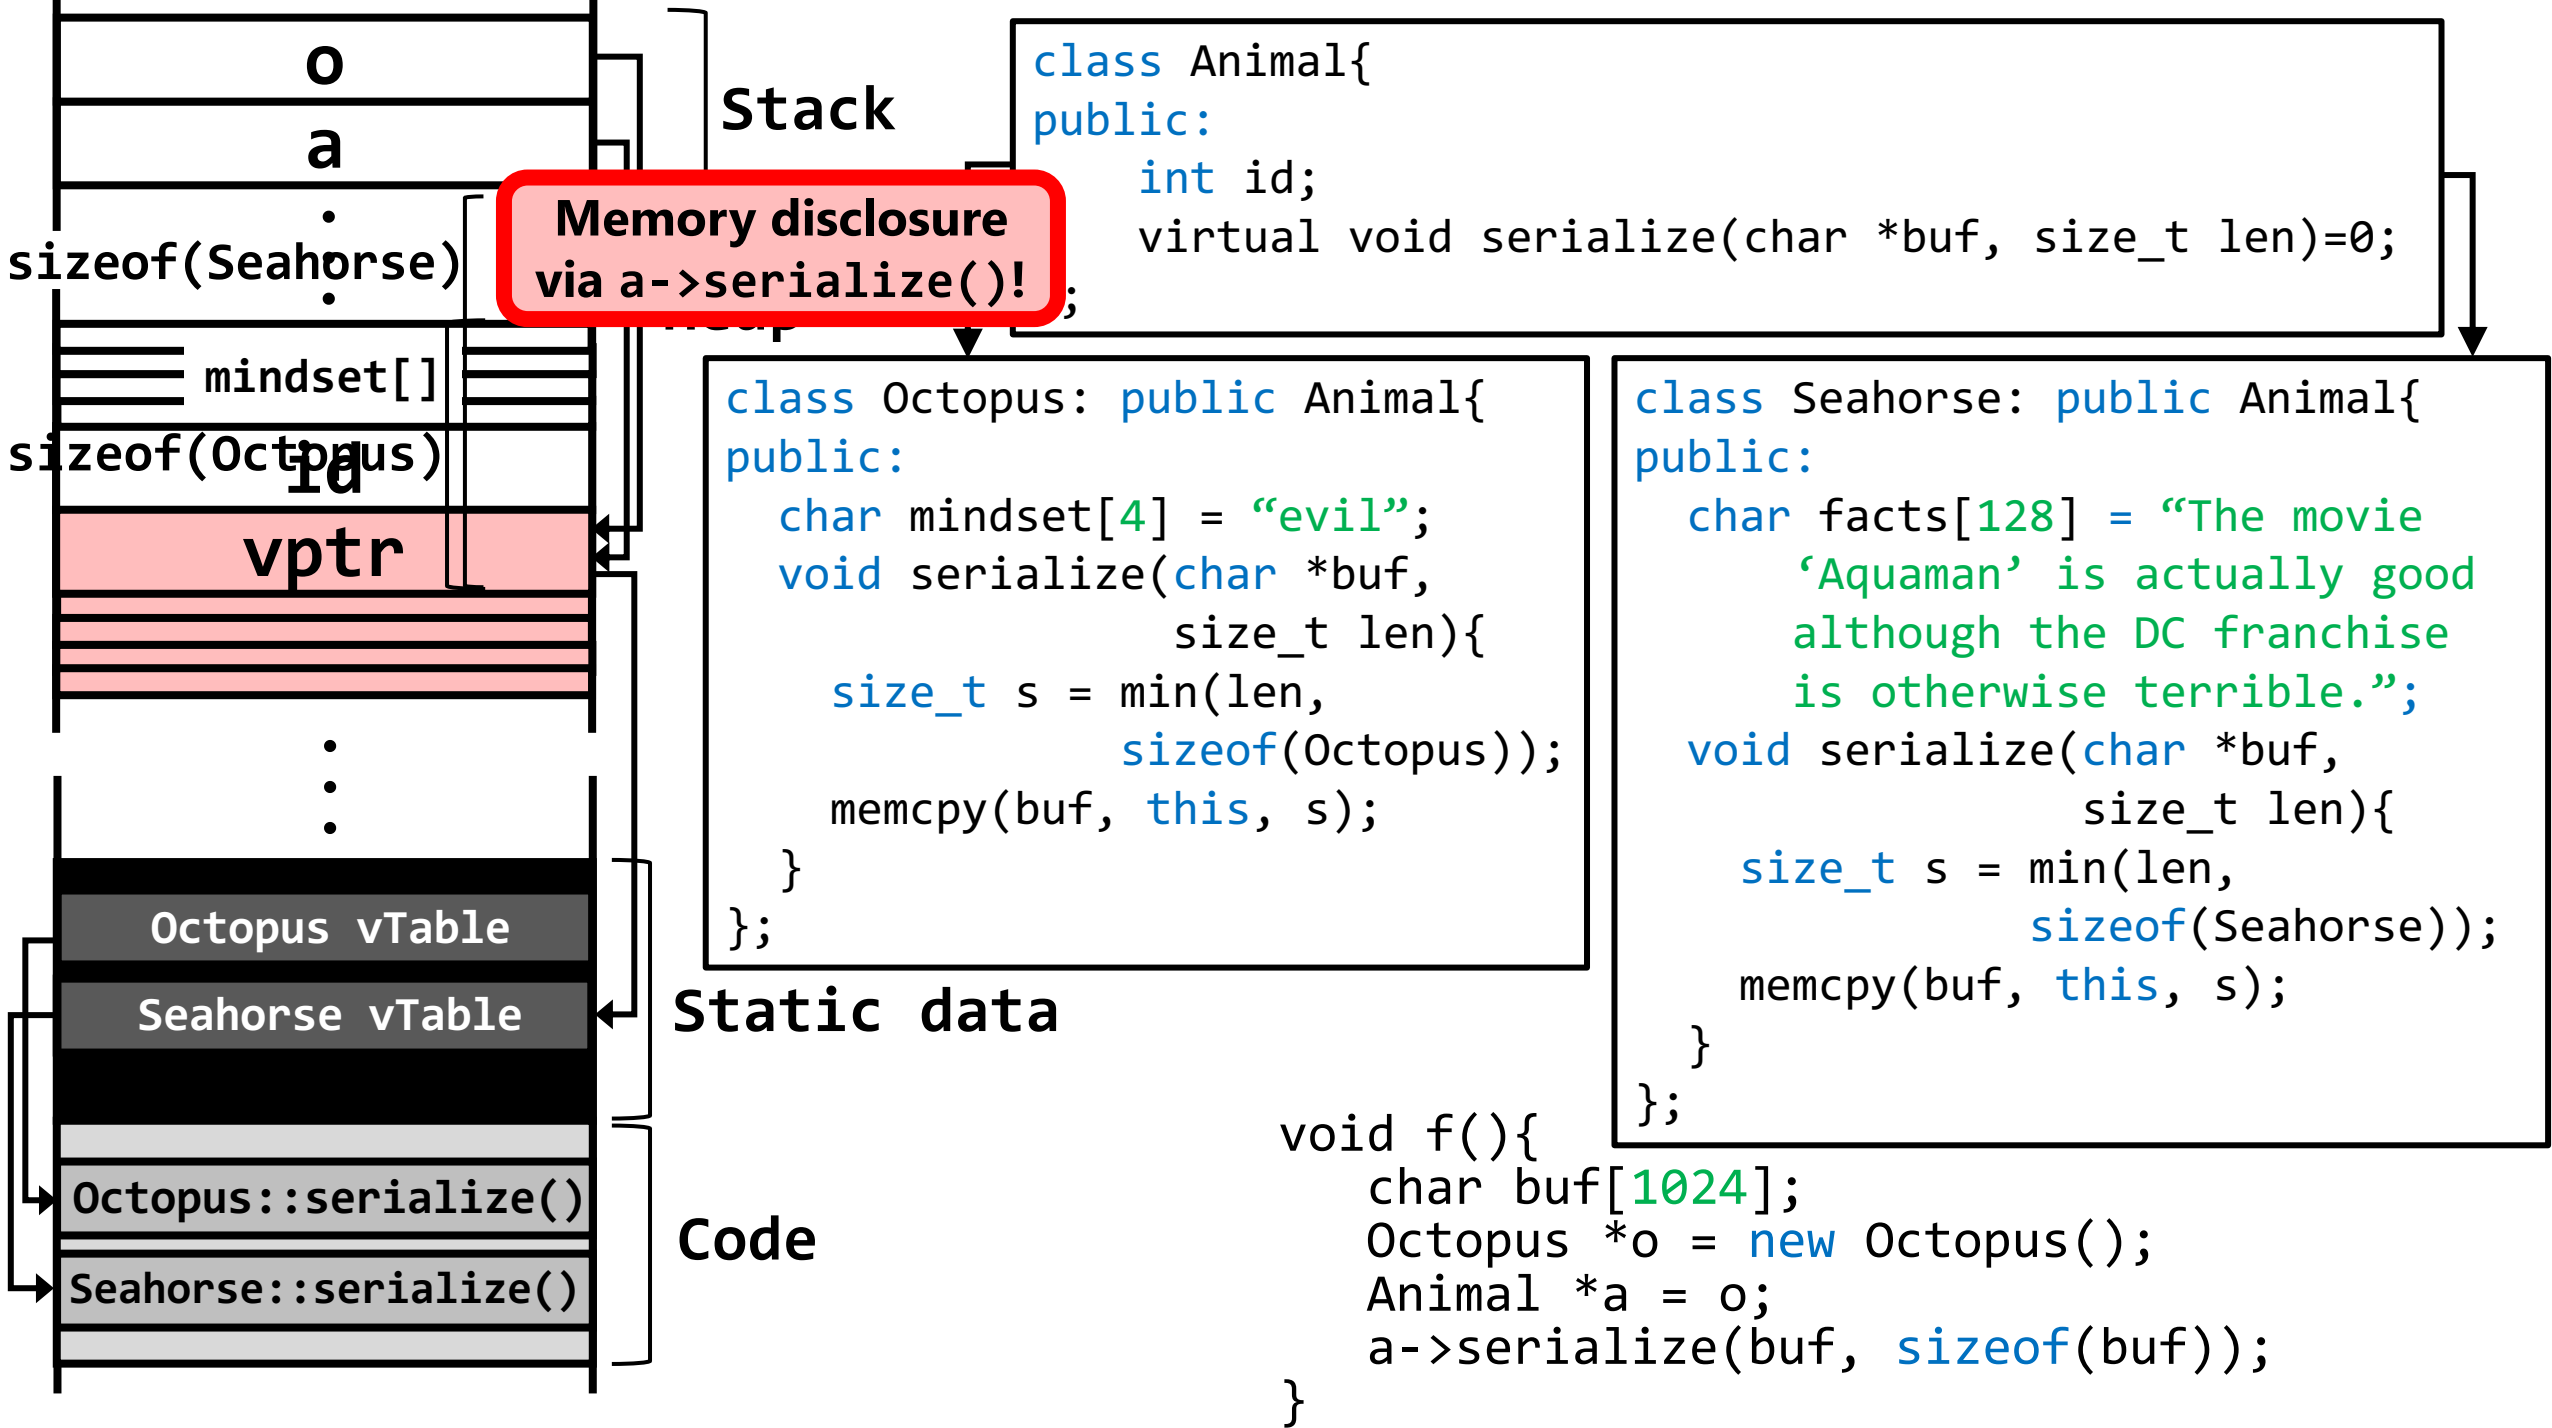
**The Octopus object invokes
the Seahorse method!**



**TYPE
CONFUSION
IS VERY,
VERY BAD**

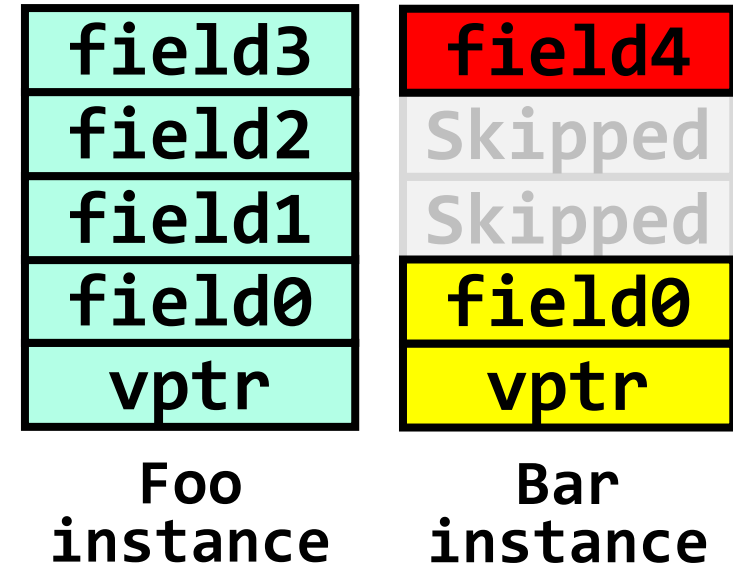






Type Confusions and Memory Read/Write Primitives

- Attackers often use type confusions to generate out-of-bounds memory accesses
 - The previous slide demonstrated a memory read vulnerability
 - Type confusion can also generate memory write vulnerabilities (e.g., assigning to an object field via a type-confused **this** pointer)
- Type confusions are very helpful if they enable non-contiguous read and write vulnerabilities
 - Helps writes avoid canaries
 - Reduces amount of unnecessary data read

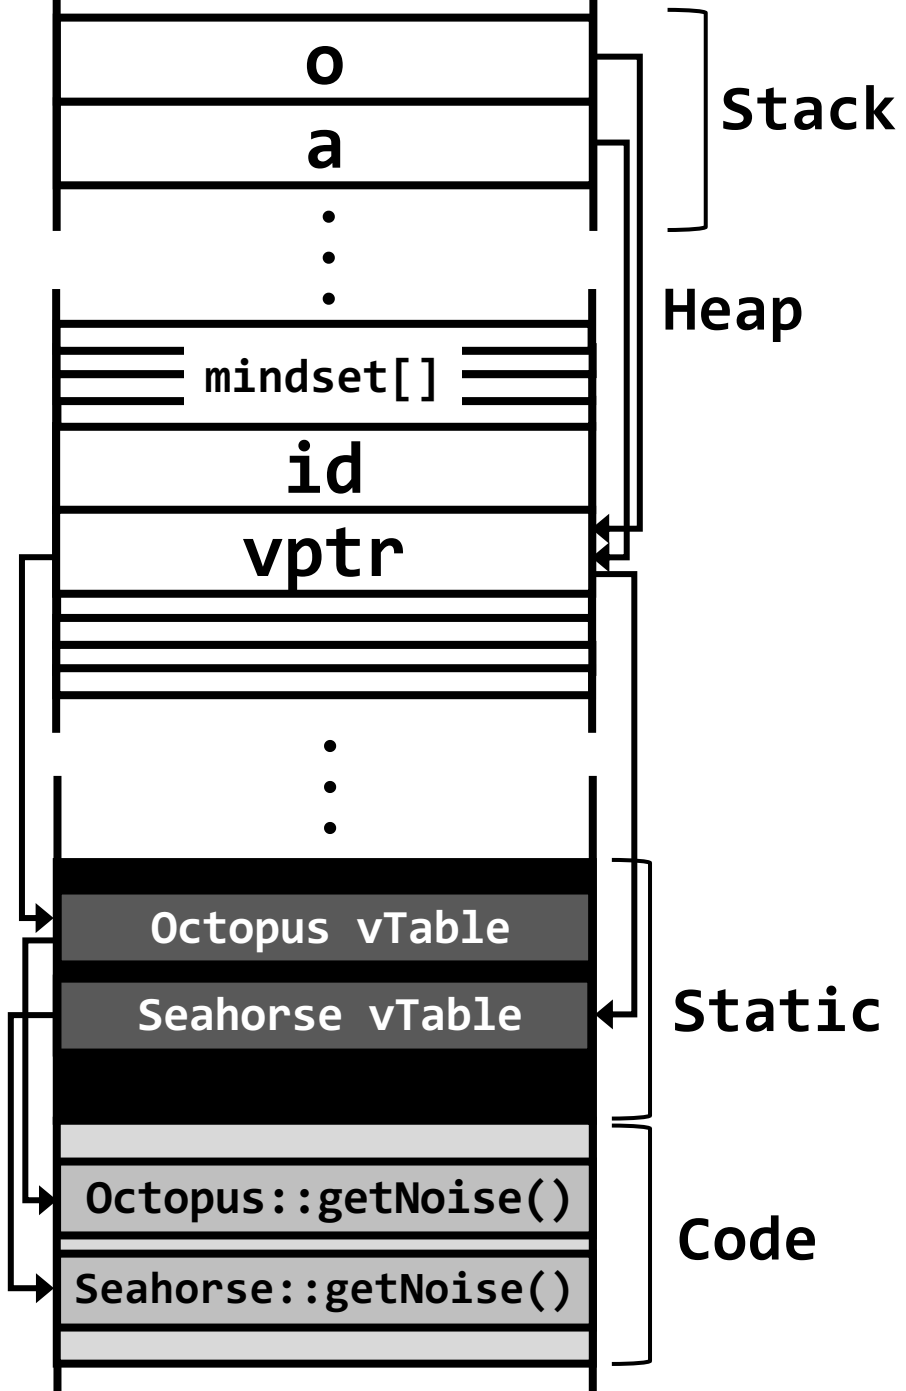


```
void Foo::set_f3(int v){  
    this->field3 = v;  
}
```

```
Bar *b = new Bar();  
Foo *f = typeConfuse(b);  
    //Exploit: f actually  
    //points to a Bar!  
f->set_f3(42); //This is  
    //a non-contiguous  
    //write vuln!
```


Type Confusions and Memory Read/Write Primitives

- Type confusions can be triggered in many ways, not just via buffer overflows!
- Ex: CVE-2015-0336
 - Malicious Flash code uses a carefully-crafted **__proto__** inheritance chain to trick a browser into interpreting a **Vector<uint>** as a **NetConnection**
 - By calling a **NetConnection** method on a **Vector<uint>**, the attacker overwrites the vector's **.length** field to be a huge value, granting read/write access to a huge swath of memory!



Overflowing a heap object




How does the attacker position the overflowable buffer beneath the vulnerable object?

Heap Feng Shui

- Suppose the attacker knows details about the heap allocation algorithm
 - Maybe it's deterministic?
 - Maybe it uses slab allocation?
 - Maybe it allocates large objects on page-aligned offsets, e.g., because `mmap()` is used to allocate the memory?
- Attacker can leverage this knowledge to place a vulnerable object next to an overflowable one!


Ex: A deterministic slab allocator



Step 1: No objects have been allocated



Step 2: The attacker allocates a bunch of objects, filling the slab




Step 3: The attacker deallocates two objects.

Heap overflow!



Step 4: The attacker allocates the overflowable object, then the vulnerable one!

Return-oriented Programming



NX prevents an attacker from injecting new code that does evil . . .

. . .but what if the attacker can create evil by chaining together snippets of preexisting code?



NAME [top](#)

gets - get a string from standard input (DEPRECATED)

SYNOPSIS [top](#)

```
#include <stdio.h>
```

```
char *gets(char *s);
```

**DESCRIPTION** [top](#)

Never use this function.

gets() reads a line from *stdin* into the buffer pointed to by *s* until either a terminating newline or **EOF**, which it replaces with a null byte ('\0'). No check for buffer overrun is performed (see BUGS below).

Suppose that a machine uses NX, but does not use stack canaries or ASLR, and contains this code . . .

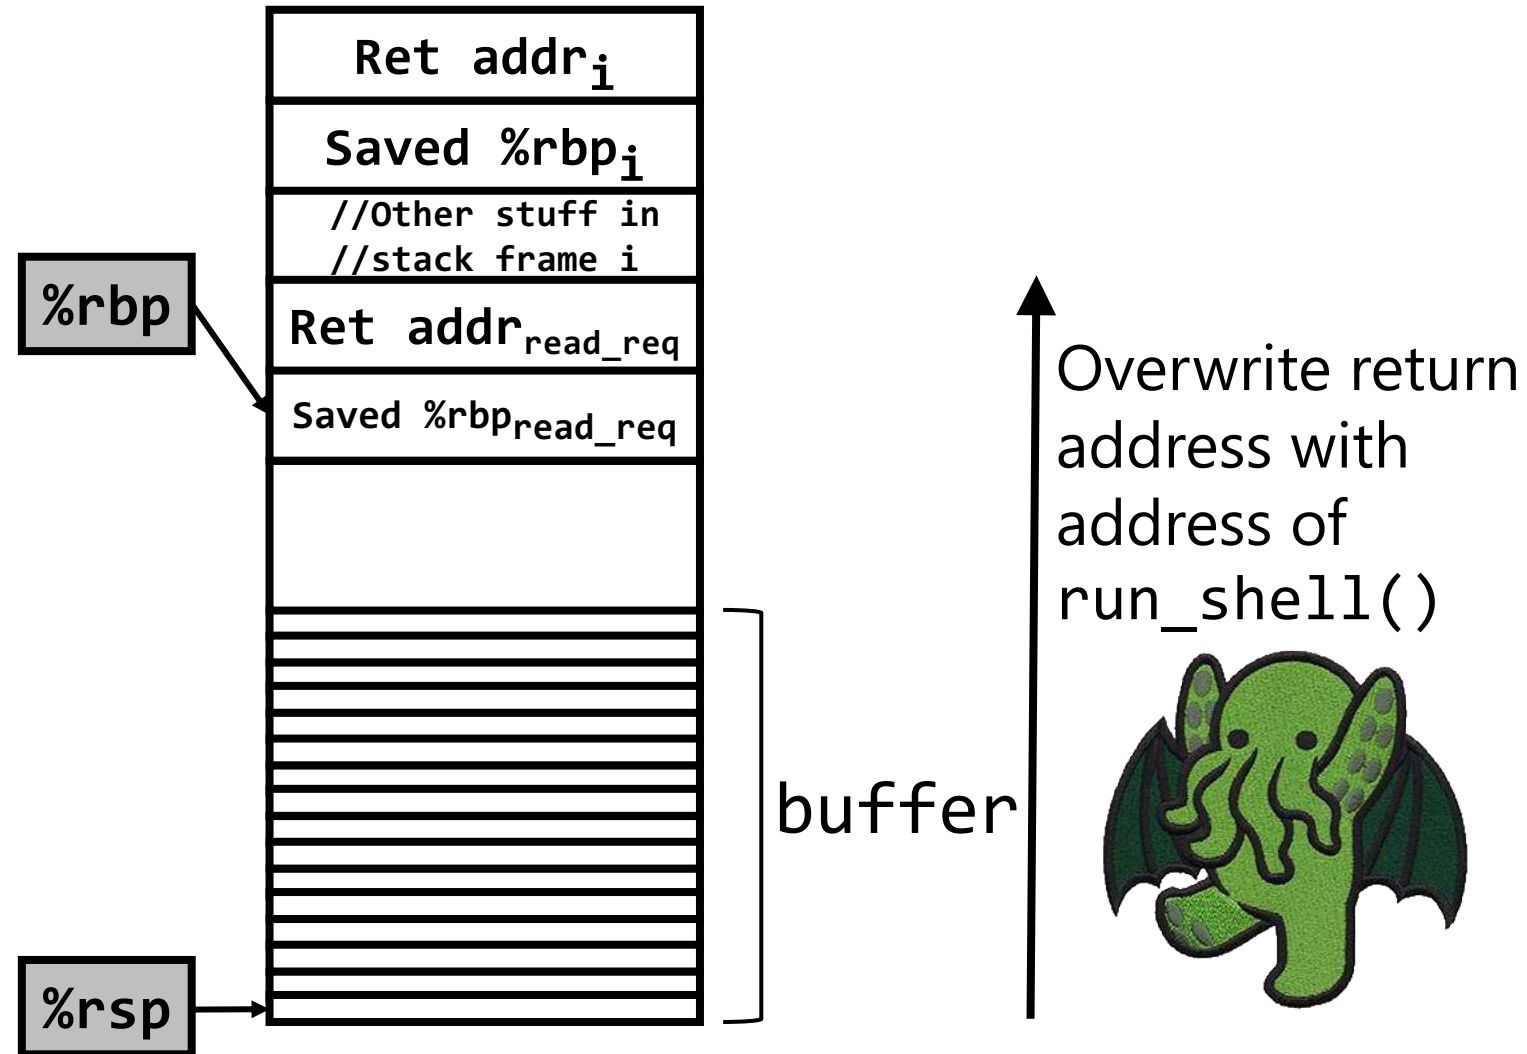
```
void read_req(){  
    char buffer[128];  
    gets(buf);  
}
```

```
void run_shell(){  
    system("/bin/sh");  
    //Spawn a new process with fork(), then do  
    //  execl("/bin/sh", "sh", "-c", command, NULL)  
    //(where the command in this example is "/bin/sh"), and wait for the  
    //command to complete; the spawned processes will use the parent's  
    //stdin/stdout/stdin due to fork()+exec() semantics.  
}
```

Suppose that a machine uses NX, but does not use stack canaries or ASLR, and contains this code . . .

```
void read_req(){
    char buffer[128];
    gets(buf);
}

void run_shell(){
    system("/bin/sh");
}
```

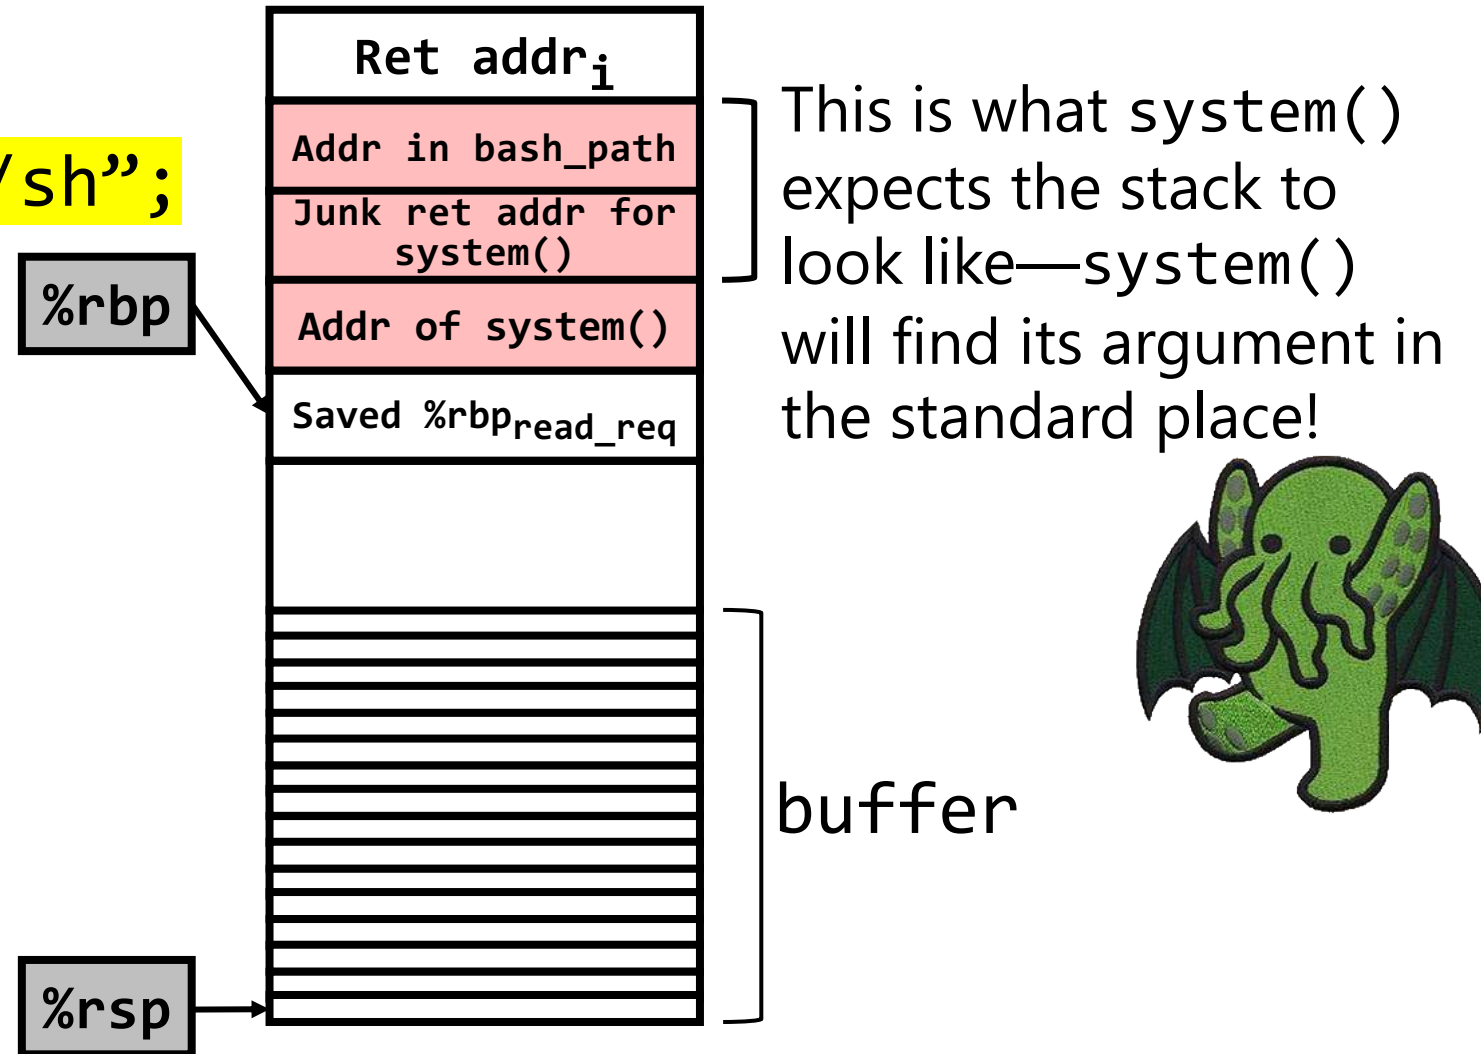


Suppose that a machine uses NX, but does not use stack canaries or ASLR, and contains this code . . .

```
char *bash_path = "/bin/sh";
```

```
void read_req(){  
    char buffer[128];  
    gets(buf);  
}
```

```
void run_ls(){  
    system("/bin/ls");  
}
```



What if the string `"/bin/sh"` isn't in the program at all?



Maybe attacks are impossible!



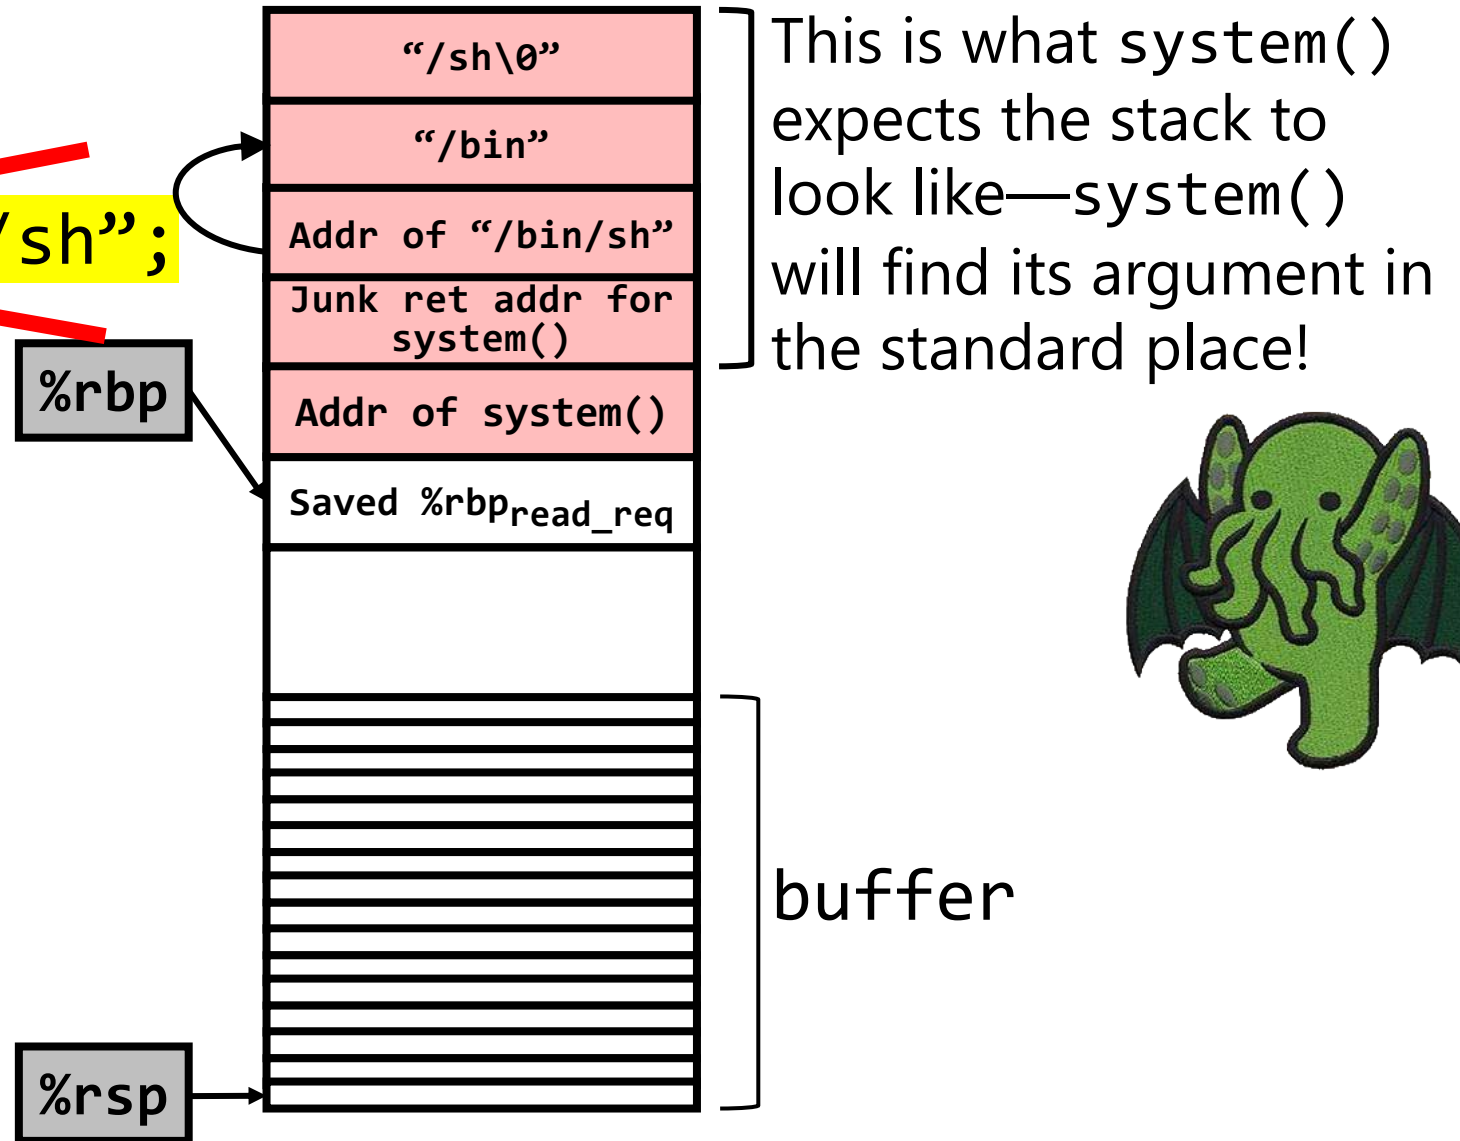
**FALSE
ATTACKS ARE VERY POSSIBLE**

Suppose that a machine uses NX, but does not use stack canaries or ASLR, and contains this code . . .

```
char *bash_path = "/bin/sh";
```

```
void read_req(){  
    char buffer[128];  
    gets(buf);  
}
```

```
void run_ls(){  
    system("/bin/ls");  
}
```



Attacker goal: Call `system("/bin/sh")` two times

- Assume that the attacker knows three addresses:
 1. The address of `system()`
 2. The address of the string `"/bin/sh"` (remember that, if necessary, the attacker can push the string onto the stack!)
 3. The starting address for this block of opcodes:

```
pop %eax    //Pops top-of-stack into %eax
ret         //Pops top-of-stack and puts
           //it in %eip
```

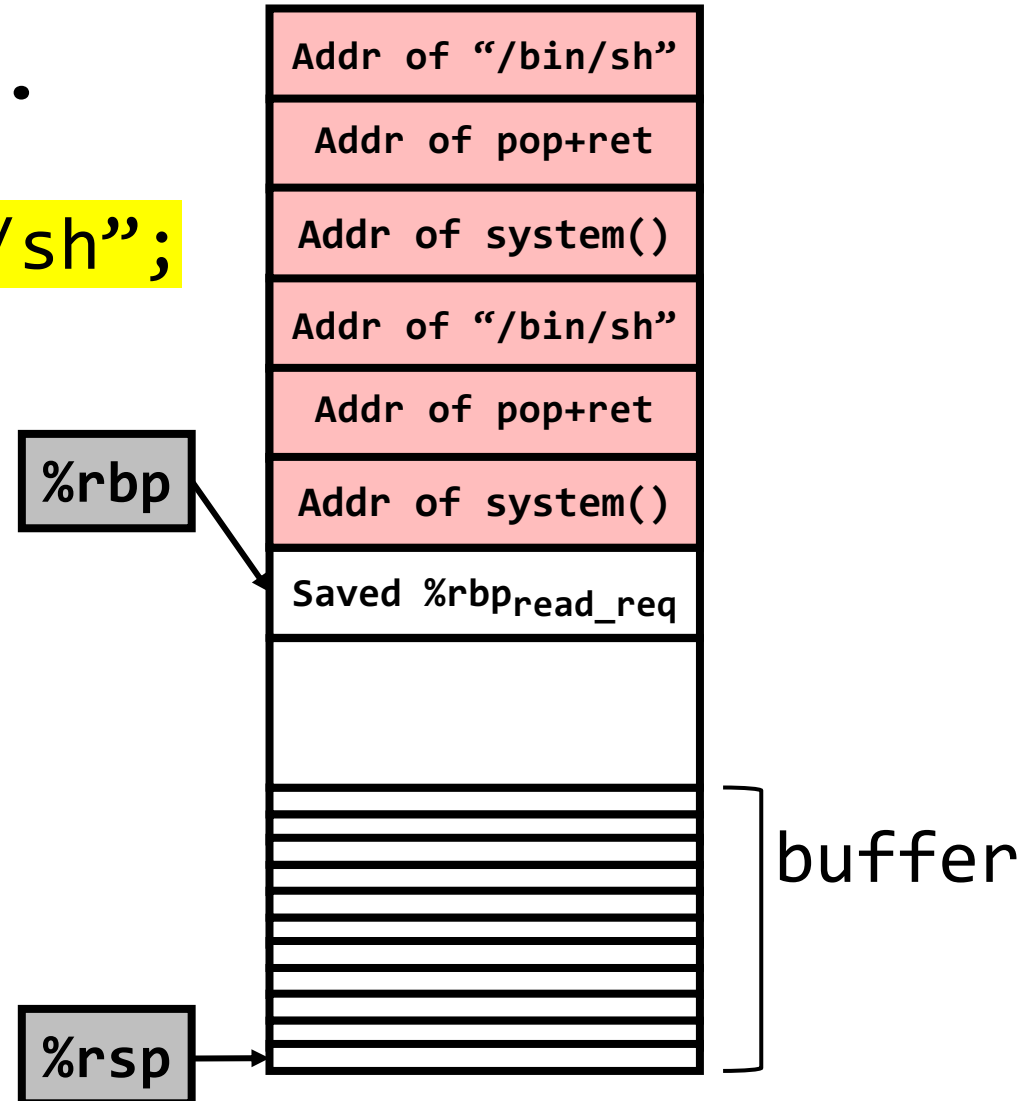
A block of opcodes is a "gadget." [There are user-friendly tools to discover gadgets in a binary, e.g., **ROPgadget**.]

Suppose that a machine uses NX, but does not use stack canaries or ASLR, and contains this code . . .

```
char *bash_path = "/bin/sh";
```

```
void read_req(){  
    char buffer[128];  
    gets(buf);  
}
```

```
void run_ls(){  
    system("/bin/ls");  
}
```



Suppose that a machine uses NX, but does not use stack canaries or ASLR, and contains this code . . .

```
char *bash_path = "/bin/sh";
```

```
void read_req(){  
    char buffer[128];  
    gets(buf);  
}
```

```
void run_ls(){  
    system("/bin/ls");  
}
```

Addr of "/bin/sh"
Addr of pop+ret
Addr of system()
Addr of "/bin/sh"
Addr of pop+ret
Addr of system()

The ret in the gadget is the pop+ret instruction. Right before read_req() invokes the gadget . . . instruction in system().

Non-trivial programs have a Turing-complete set of gadgets!

JonathanSalwan / ROPgadget

```
- Step 4 -- Syscall gadget
```

```
[+] Gadget found: 0x804936d int 0x80
```

```
[+] Gadget found: 0x8056c56 pop ecx ; pop ebx ; ret  
[+] Gadget found: 0x8056c2c pop edx ; ret
```

SweetVishnya committed ce1ba02 1 hour ago 485 commits

.github/workflows	Update main.yml
ropgadget	Remove unused sqlite3 import
scripts	Change LICENSE: From GPL to BSD
test-suite-binaries	Fix tests

ROP chain generation

```
- Step 1 -- Write-what-where gadgets
```

```
[+] Gadget found: 0x806f702 mov dword ptr [edx], ecx ; ret  
[+] Gadget found: 0x8056c2c pop edx ; ret  
[+] Gadget found: 0x8056c56 pop ecx ; pop ebx ; ret  
[-] Can't find the 'xor ecx, ecx' gadget. Try with another 'mov [r], r'  
  
[+] Gadget found: 0x808fe0d mov dword ptr [edx], eax ; ret  
[+] Gadget found: 0x8056c2c pop edx ; ret  
[+] Gadget found: 0x80c5126 pop eax ; ret
```

```
- Step 4 -- Syscall gadget
```

```
[+] Gadget found: 0x804936d int 0x80
```

```
- Step 5 -- Build the ROP chain
```

```
#!/usr/bin/env python2  
# execve generated by ROPgadget v5.2  
  
from struct import pack  
  
# Padding goes here  
p = ''  
  
p += pack('<I', 0x08056c2c) # pop edx ; ret  
p += pack('<I', 0x080f4060) # @ .data  
p += pack('<I', 0x080c5126) # pop eax ; ret  
p += '/bin'
```


Instruction Prefixes	Opcode	ModR/M	SIB	Displacement	Immediate
----------------------	--------	--------	-----	--------------	-----------

Prefixes of
1 byte each
(if required)

1-, 2-, or 3-byte
opcode

1 byte
(if required)


1 byte
(if required)

Address

Immediate

**x86: The Oil Isn't
Washing Off**






Used by 32-bit Linux processes to invoke a system call!

and %eax, 0x000080cd


Gadget 2



mov %edx,0x4(%eax)

mov %eax,%ds:0x0805d0ff

Gadget 3



add 0x4,%esp

pop %ebx

pop %ebp ret

Gadget 1

Jumping into this range will cause the CPU to raise an "invalid opcode" exception!



- In ROP, attacker creates a program whose execution is driven by the stack pointer!
- As the stack pointer moves down the stack (i.e., UP in memory), the program executes gadgets whose code is preexisting in the executable
- This attack can't be stopped by NX bits!
- But what about stack canaries?

Defeating Stack Canaries



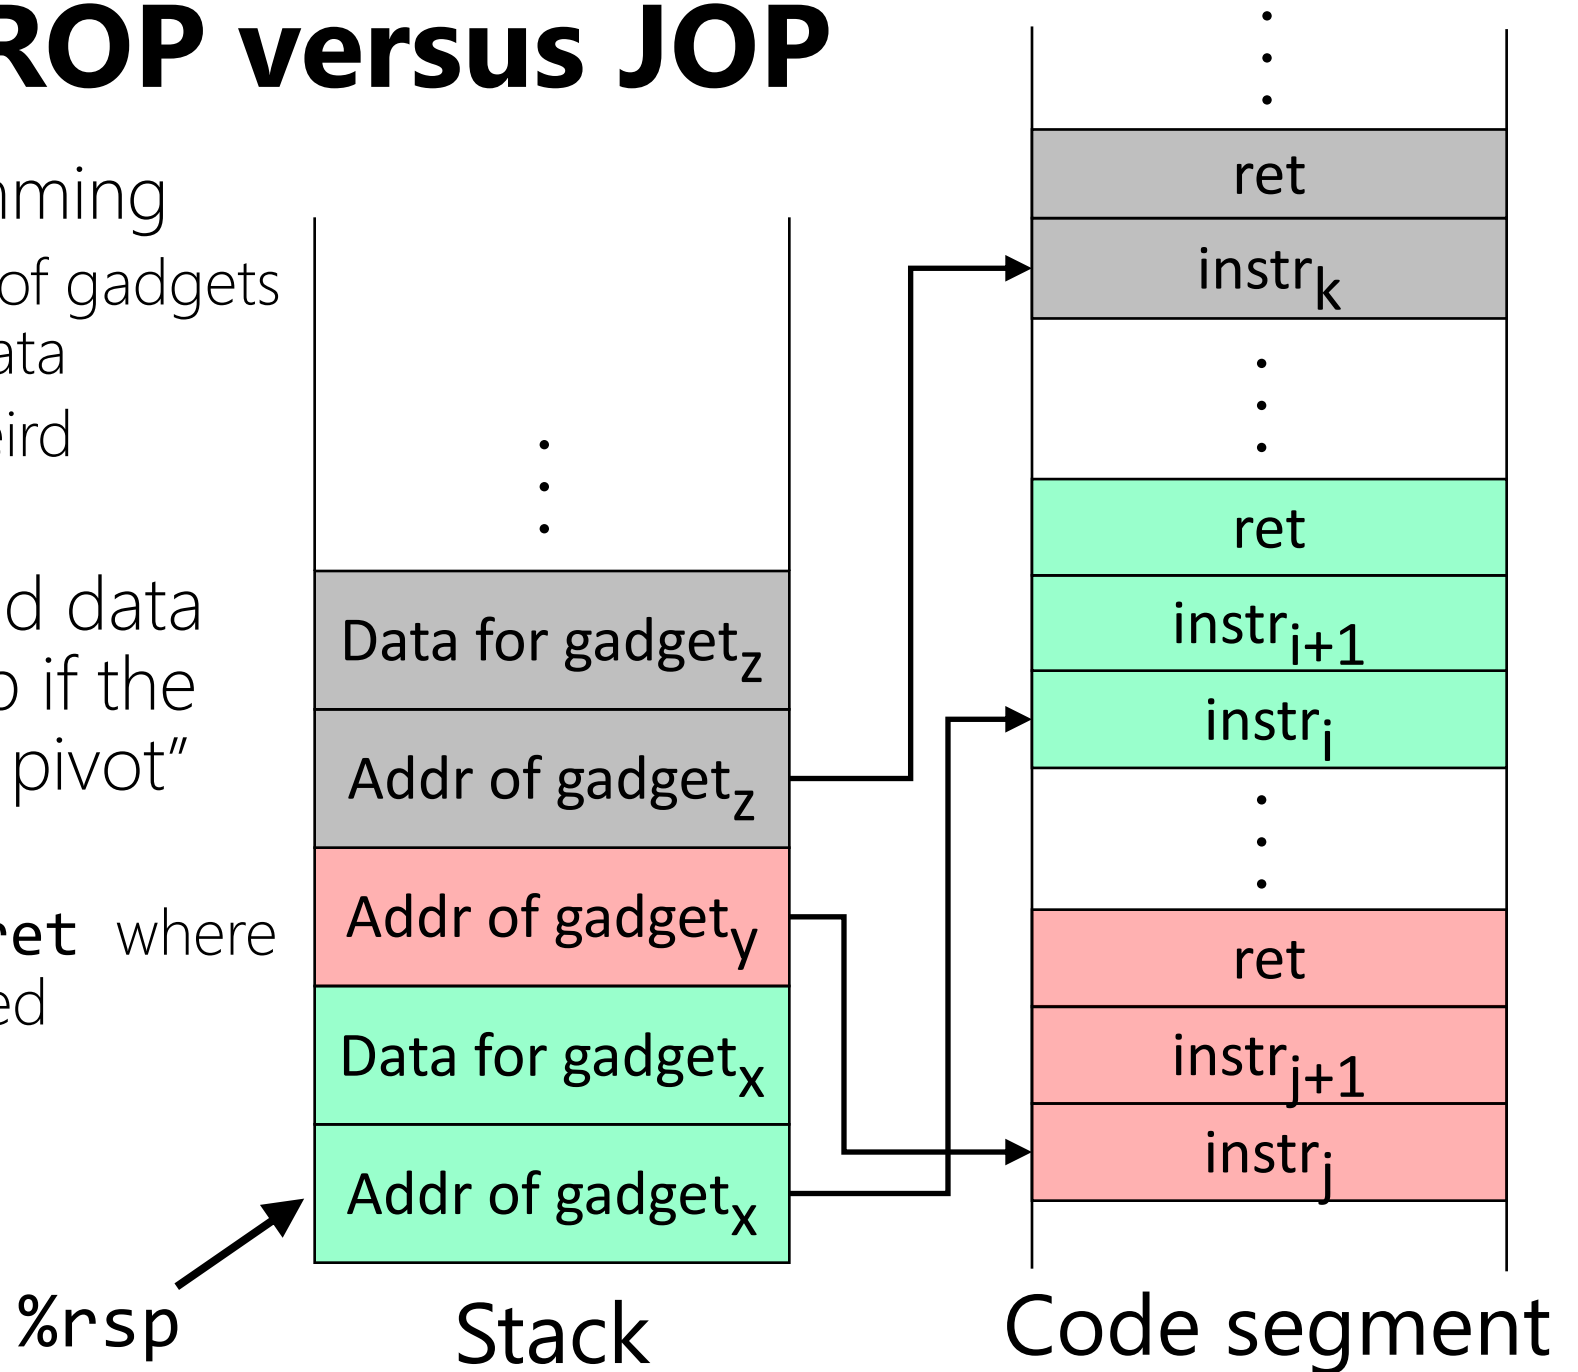
- Approach 1: Non-contiguous writes
 - A non-contiguous write vulnerability allows the attacker to only write locations of interest (but the attacker must know where to write . . .)
 - Type confusions are a common source of non-contiguous writes
- Approach 2: Stack pivoting
 - The attacker corrupts a register that is later used by the program to modify the stack
 - Ex: `mov %rax, %rsp` where `%rax` is attacker controlled

Defeating ASLR

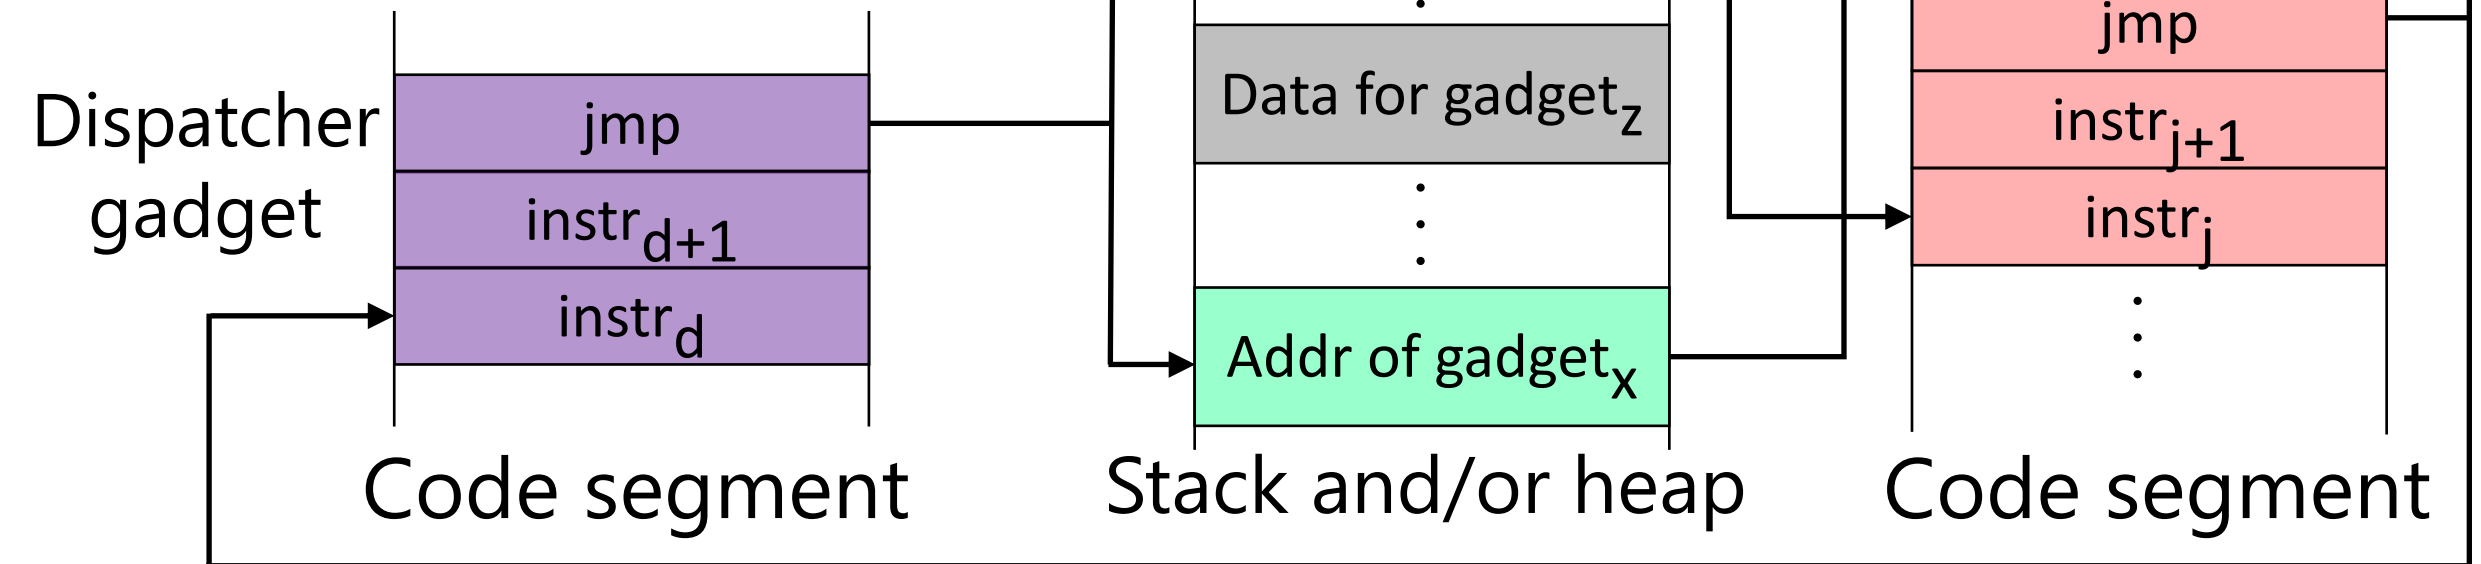
- Usually attackers try to use memory read vulnerabilities to find a single concrete memory address that enables the discovery of other concrete addresses
 - Ex: Discovering a single code pointer to a function in **libc** lets the attacker determine the location of any function in **libc** (since those relative offsets aren't changed by ASLR!)
 - Ex: An over-read of a buffer on the stack can disclose return addresses, narrowing the location of the stack
- Attackers can also leverage program quirks
 - Ex: A program might load non-relocatable libraries—look for gadgets there!
 - Ex: ASLR doesn't change after **fork()**, so the attacker can repeatedly guess locations and reuse work if crashes occur and new processes are respawned

ROP versus JOP

- Return-oriented programming
 - Stack hold the addresses of gadgets as well as gadget input data
 - Stack pointer acts as a weird program counter
- The gadget addresses and data can be stored in the heap if the attacker can find a “stack pivot” gadget
 - Ex: `mov %rax, %rsp; ret` where `%rax` is attacker-controlled




- Jump-oriented programming
 - A dispatch table holds addresses of gadgets as well as gadget input data
 - A special “dispatch gadget” increments the program counter
- The dispatch table doesn't need to be in contiguous memory!
 - In contrast, ROP attack requires gadget addresses + data to be contiguous



A real dispatcher gadget in libc

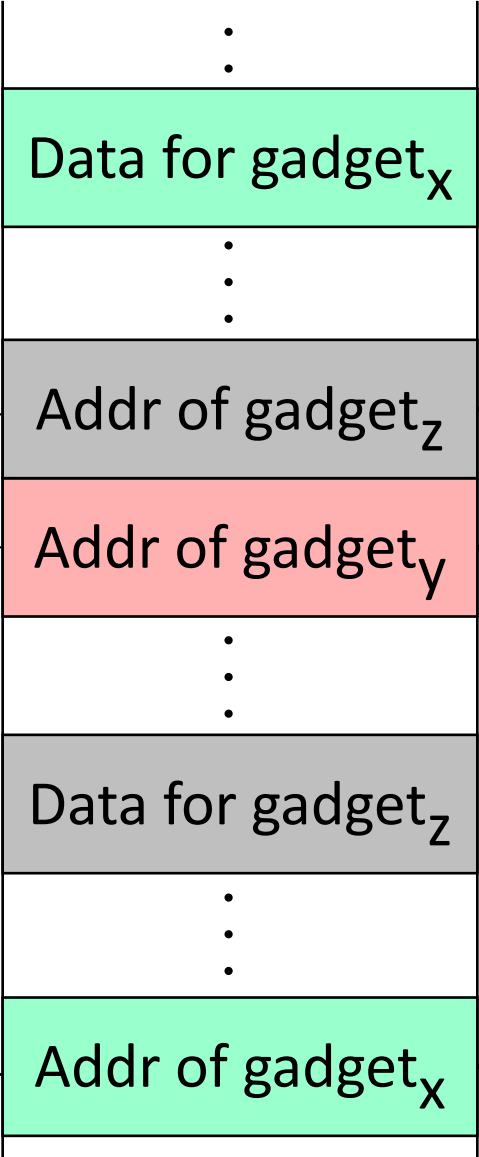
```
add %edi, %ebp
jmp *(%ebp - 0x39)
//Attacker corrects for 0x39
//by using prior gadgets to
//massage %edi
```

Dispatcher
gadget




Code segment

Dispatch table



Stack and/or heap



Code segment

A real dispatcher gadget in libc

```
add %edi, %ebp
jmp *(%ebp - 0x39)
//Attacker corrects for 0x39
//by using prior gadgets to
//massage %edi
```

A real JOP attack using libc gadgets

```
popa                //Load all registers from
                    //attacker-controlled stack

cmc                //No practical effect
jmp *(%ecx)         //Go to dispatcher

xchg %ecx, %eax     //Swap reg values
fdiv st3, st0       //No practical effect
jmp *(%esi - 0xf)   //Go to dispatcher

mov *(%esi + 0xc), %eax //Set %eax
mov %eax, %esp      //No practical effect
call *(%esi + 0x4)  //Go to dispatcher

sysenter            //Invoke a system call,
                    //with the argument
                    //registers %eax (the,
                    //syscall num), %ebx,
                    //ecx, %edx set using
                    //popa and gadgets
```

**THIS ALL SOUNDS TERRIBLE
BUT IT'S NOT AS BAD AS IT SEEMS**





Control-flow Enforcement Technology Specification

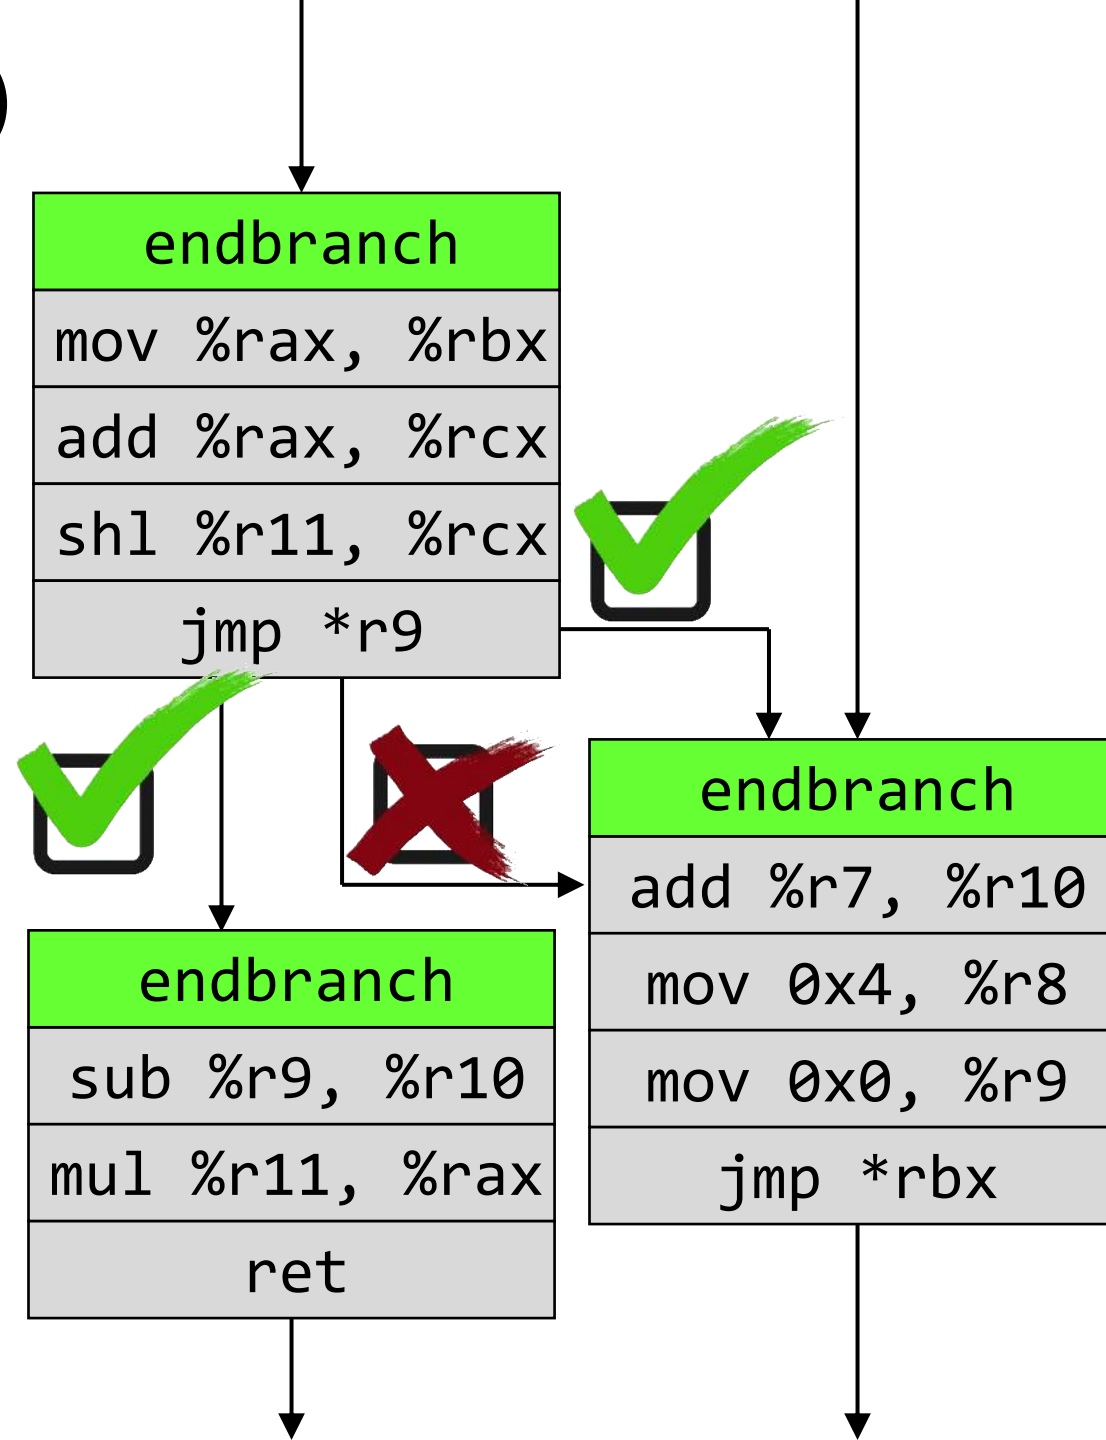
May 2019
Revision 3.0

Document Number: 334525-003

- Intel's CET spec introduces two technologies
 - Shadow stacks (as discussed last lecture!)
 - Indirect branch tracing (which we'll discuss next)
- CET ships with:
 - Tiger Lake chips that arrive this year
 - Upcoming Xeon chips
- AMD says that it's working on equivalent technology

Indirect Branch Tracking (IBT)

- IBT introduces **endbranch** (which evaluates to **nop** on older CPUs)
 - Hardware ensures that any **jmp/call** target is an **endbranch** . . .
 - . . . unless the **jmp/call** instruction is prefixed with **3EH** ("no-track")
- To use IBT:
 - Set configuration registers like:
 - "Master enable bit" **%cr4.cet**
 - "User-mode IBT enable bit" **%ia32_u_cet.endbr_en**
 - Use a compiler that inserts **endbranch** at each potential indirect jump target



`-fcf-protection=[full branch return none]` → **Enable IBT and shadow stacks**

Enable code instrumentation of control-flow transfers to increase program security by checking that target addresses of control-flow transfer instructions (such as indirect function call, function return, indirect jump) are valid. This prevents diverting the flow of control to an unexpected target. This is intended to protect against such threats as Return-oriented Programming (ROP), and similarly call/jmp-oriented programming (COP/JOP).

The value `branch` tells the compiler to implement checking of validity of control-flow transfer at the point of indirect branch instructions, i.e. `call/jmp` instructions. The value `return` implements checking of validity at the point of returning from a function. The value `full` is an alias for specifying both `branch` and `return`. The value `none` turns off instrumentation.

The macro `__CET__` is defined when `-fcf-protection` is used. The first bit of `__CET__` is set to 1 for the value `branch` and the second bit of `__CET__` is set to 1 for the `return`.

You can also use the `nocf_check` attribute to identify which functions and calls should be skipped from instrumentation (see [Function Attributes](#)).

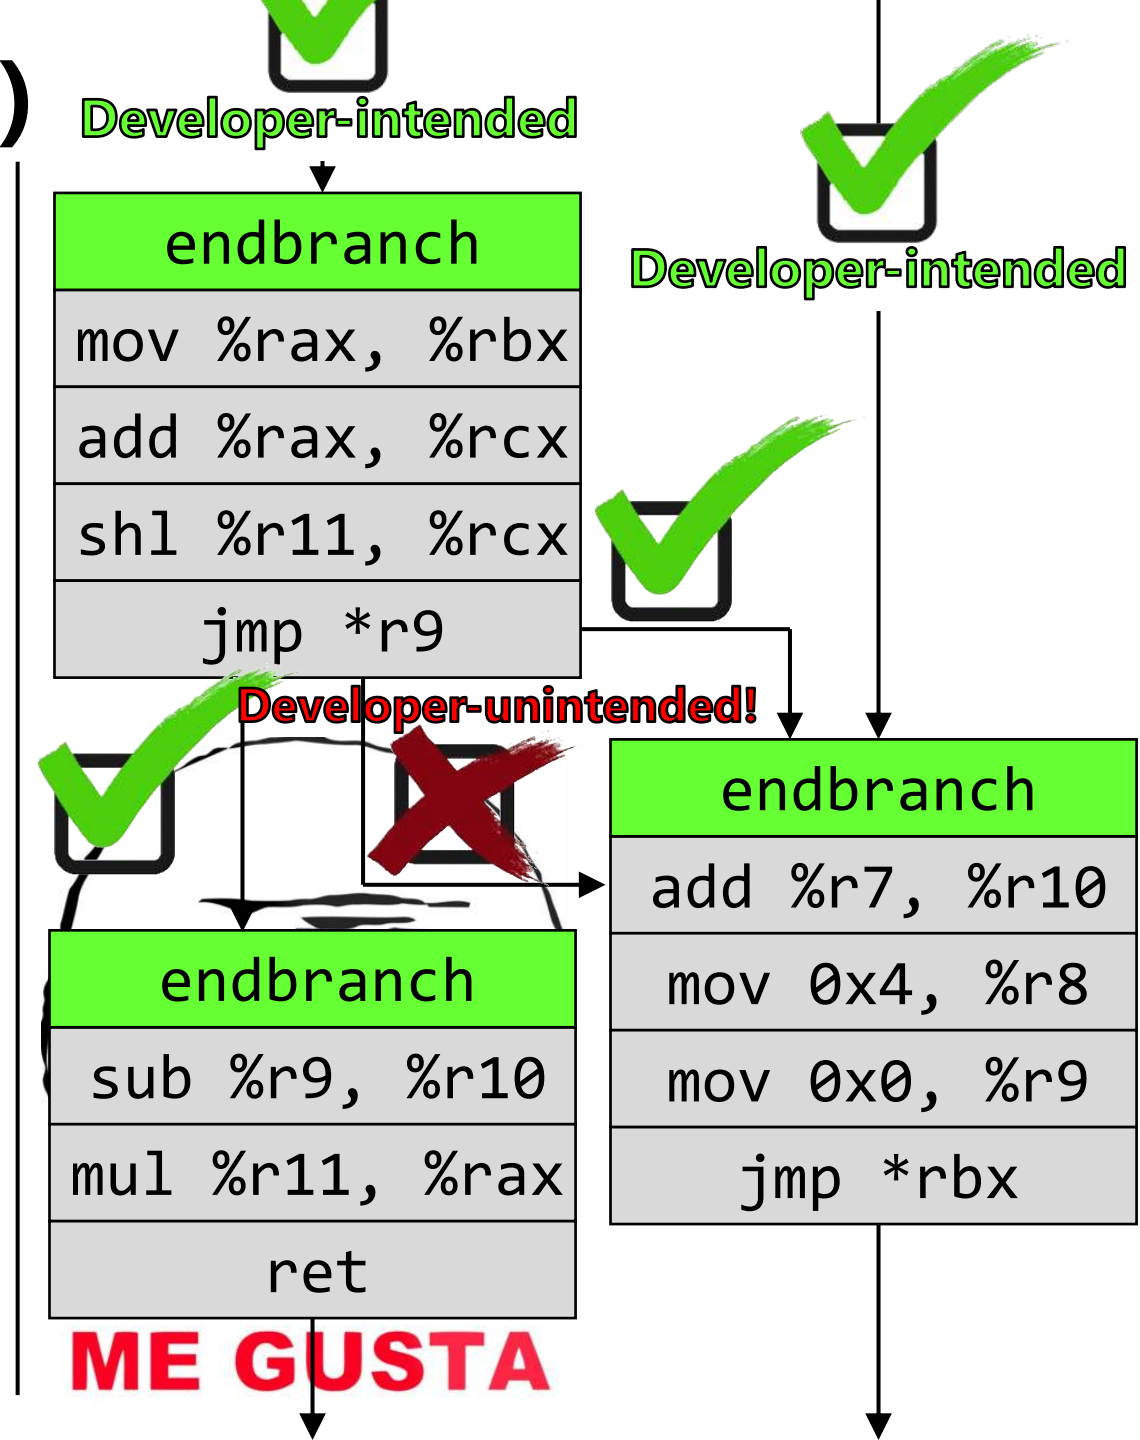
Currently the x86 GNU/Linux target provides an implementation based on Intel Control-flow Enforcement Technology (CET).

`-fstack-protector` → **Enable stack canaries**

Emit extra code to check for buffer overflows, such as stack smashing attacks. This is done by adding a guard variable to functions with vulnerable objects. This includes functions that call `alloca`, and functions with buffers larger than or equal to 8 bytes. The guards are initialized when a function is entered and then checked when the function exits. If a guard check fails, an error message is printed and the program exits. Only variables that are actually allocated on the stack are considered, optimized away variables or variables allocated in registers don't count.

Indirect Branch Tracking (IBT)

- IBT provides relaxed CFI: a given `jmp/ call` can go to any valid `jmp/call` target
- ROP and JOP are still possible if the attacker can find enough **endbranch**-preceded gadgets
- However, IBT in concert with shadow stacks, ASLR, NX bits, and stack canaries make it much harder for attackers to divert control flow!





What Do Threats Look Like Today?

Trends, challenges, and strategic shifts in the software vulnerability mitigation landscape


Matt Miller (@epakskape)

Microsoft Security Response Center (MSRC)


BlueHat IL

February 7th, 2019

More vulnerabilities fixed, fewer known exploits




On the surface, risk appears to be *increasing*



But known actualized risk appears to be *decreasing*


Memory safety issues remain dominant

We closely study the root cause trends of vulnerabilities & search for patterns



~70% of the vulnerabilities addressed through a security update each year continue to be memory safety issues

Drilling down into root causes



Stack corruptions are essentially dead

Use after free spiked in 2013-2015 due to web browser UAF, but was mitigated by Mem GC

Heap out-of-bounds read, type confusion, & uninitialized use have generally increased

Spatial safety remains the most common vulnerability category (heap out-of-bounds read/write)

Top root causes since 2016:

#1: heap out-of-bounds

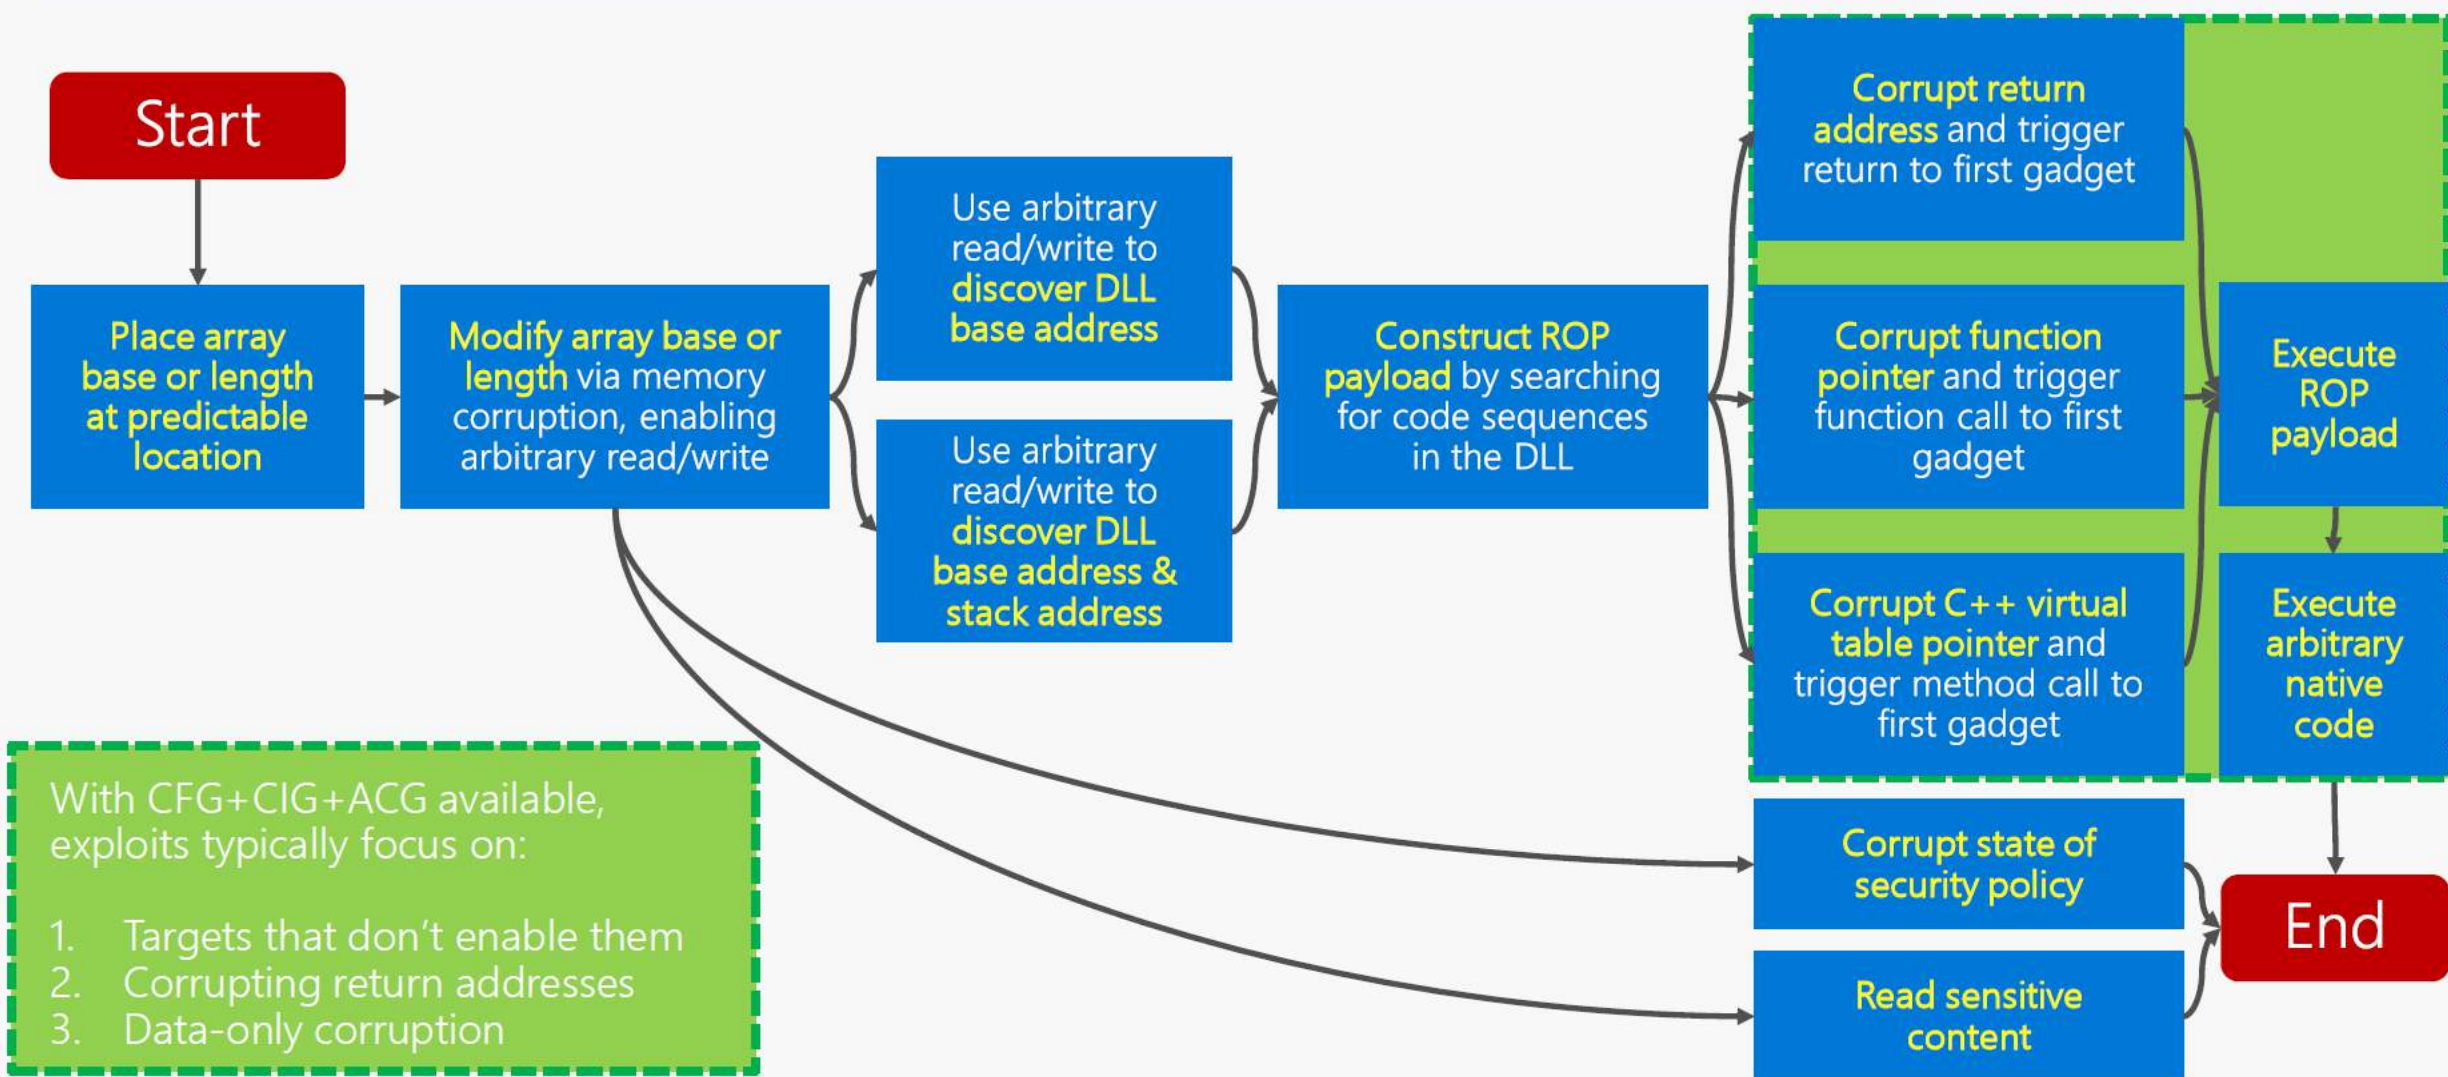
#2: use after free

#3: type confusion


#4: uninitialized use

Challenges with breaking exploitation techniques [2/4]

Most exploits have followed the same general steps since ~2016



SGX



Overview of SGX

- Goal 1: Allow a client to interact with a secure server-side computation space on an untrusted host

FFFFFFFF FFFFFFFF
 The OS and the host are untrusted


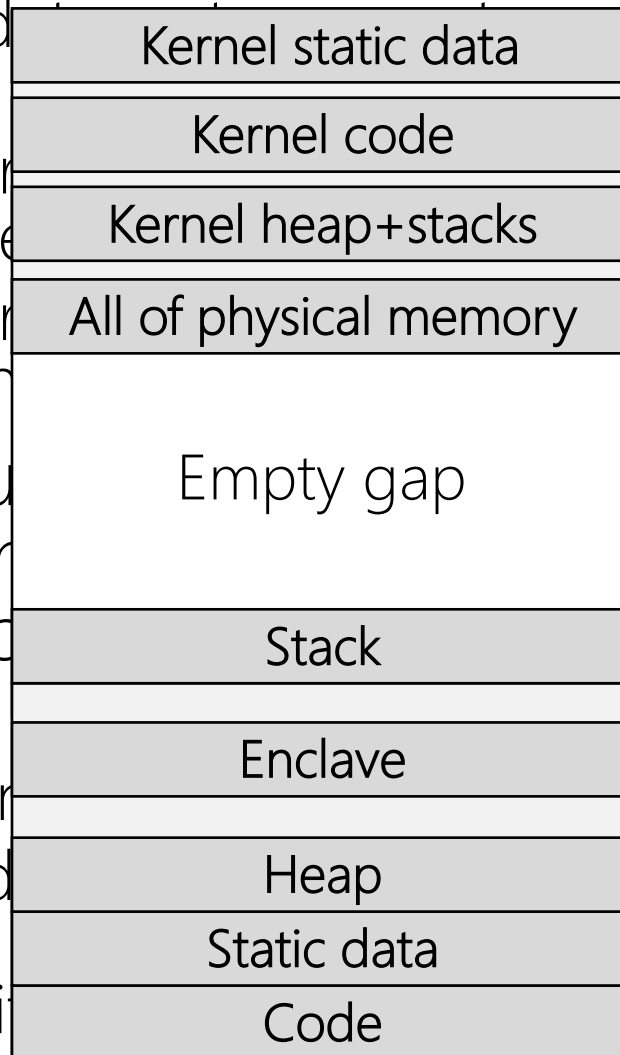
Kernel-mode

FFFFFFFF 8000 00000000
 Prevent either party from changing state without being detected

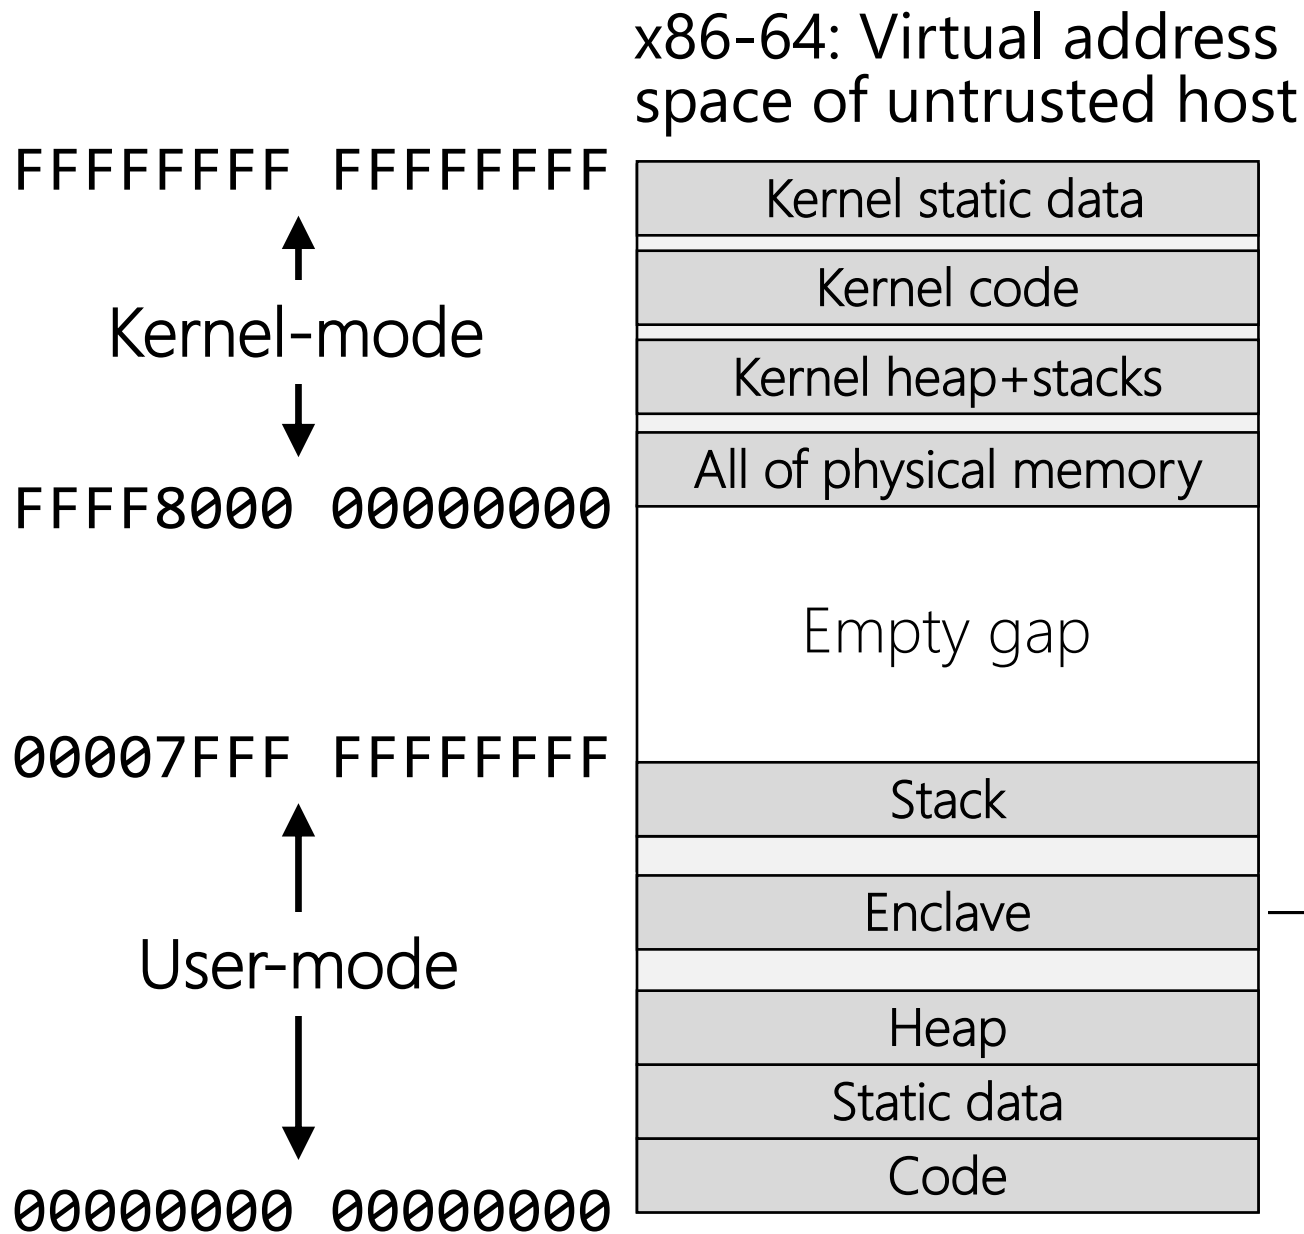
- Goal 2: Allow secure computations to execute atop the host's hardware circuits which execute untrusted computations

User-mode


00007FFF FFFFFFFF
 On no: how to add SGX hardware to legacy microarchitectures



Overview of SGX



- Untrusted host process embeds an enclave
 - Cannot access enclave pages: reads return -1, writes are ignored
 - Jumps to enclave code using **EENTER**
- Enclave code runs at Ring 3 (i.e., the least-privileged level)
 - However, can access the entire user-mode virtual address space of host
 - Cannot **syscall** or **int** instructions: relies on host for IO!
 - Data for a write IO must be placed in memory accessible to untrusted host
 - Data from a read IO must be pulled from memory accessible to untrusted host
 - Returns to untrusted host using **EEXIT** instruction



%cr3

isEnclave?=0

L1 i-cache

L1 d-cache

L2 cache

L3 cache



Memory encryption engine

TLB

virtPage

physFrame

Enclave Page Cache Map (EPCM)

Valid?

Owning enclave id
(i.e., SECS id)

Base virt. addr.
of page

Blocked?

R/W/X perms

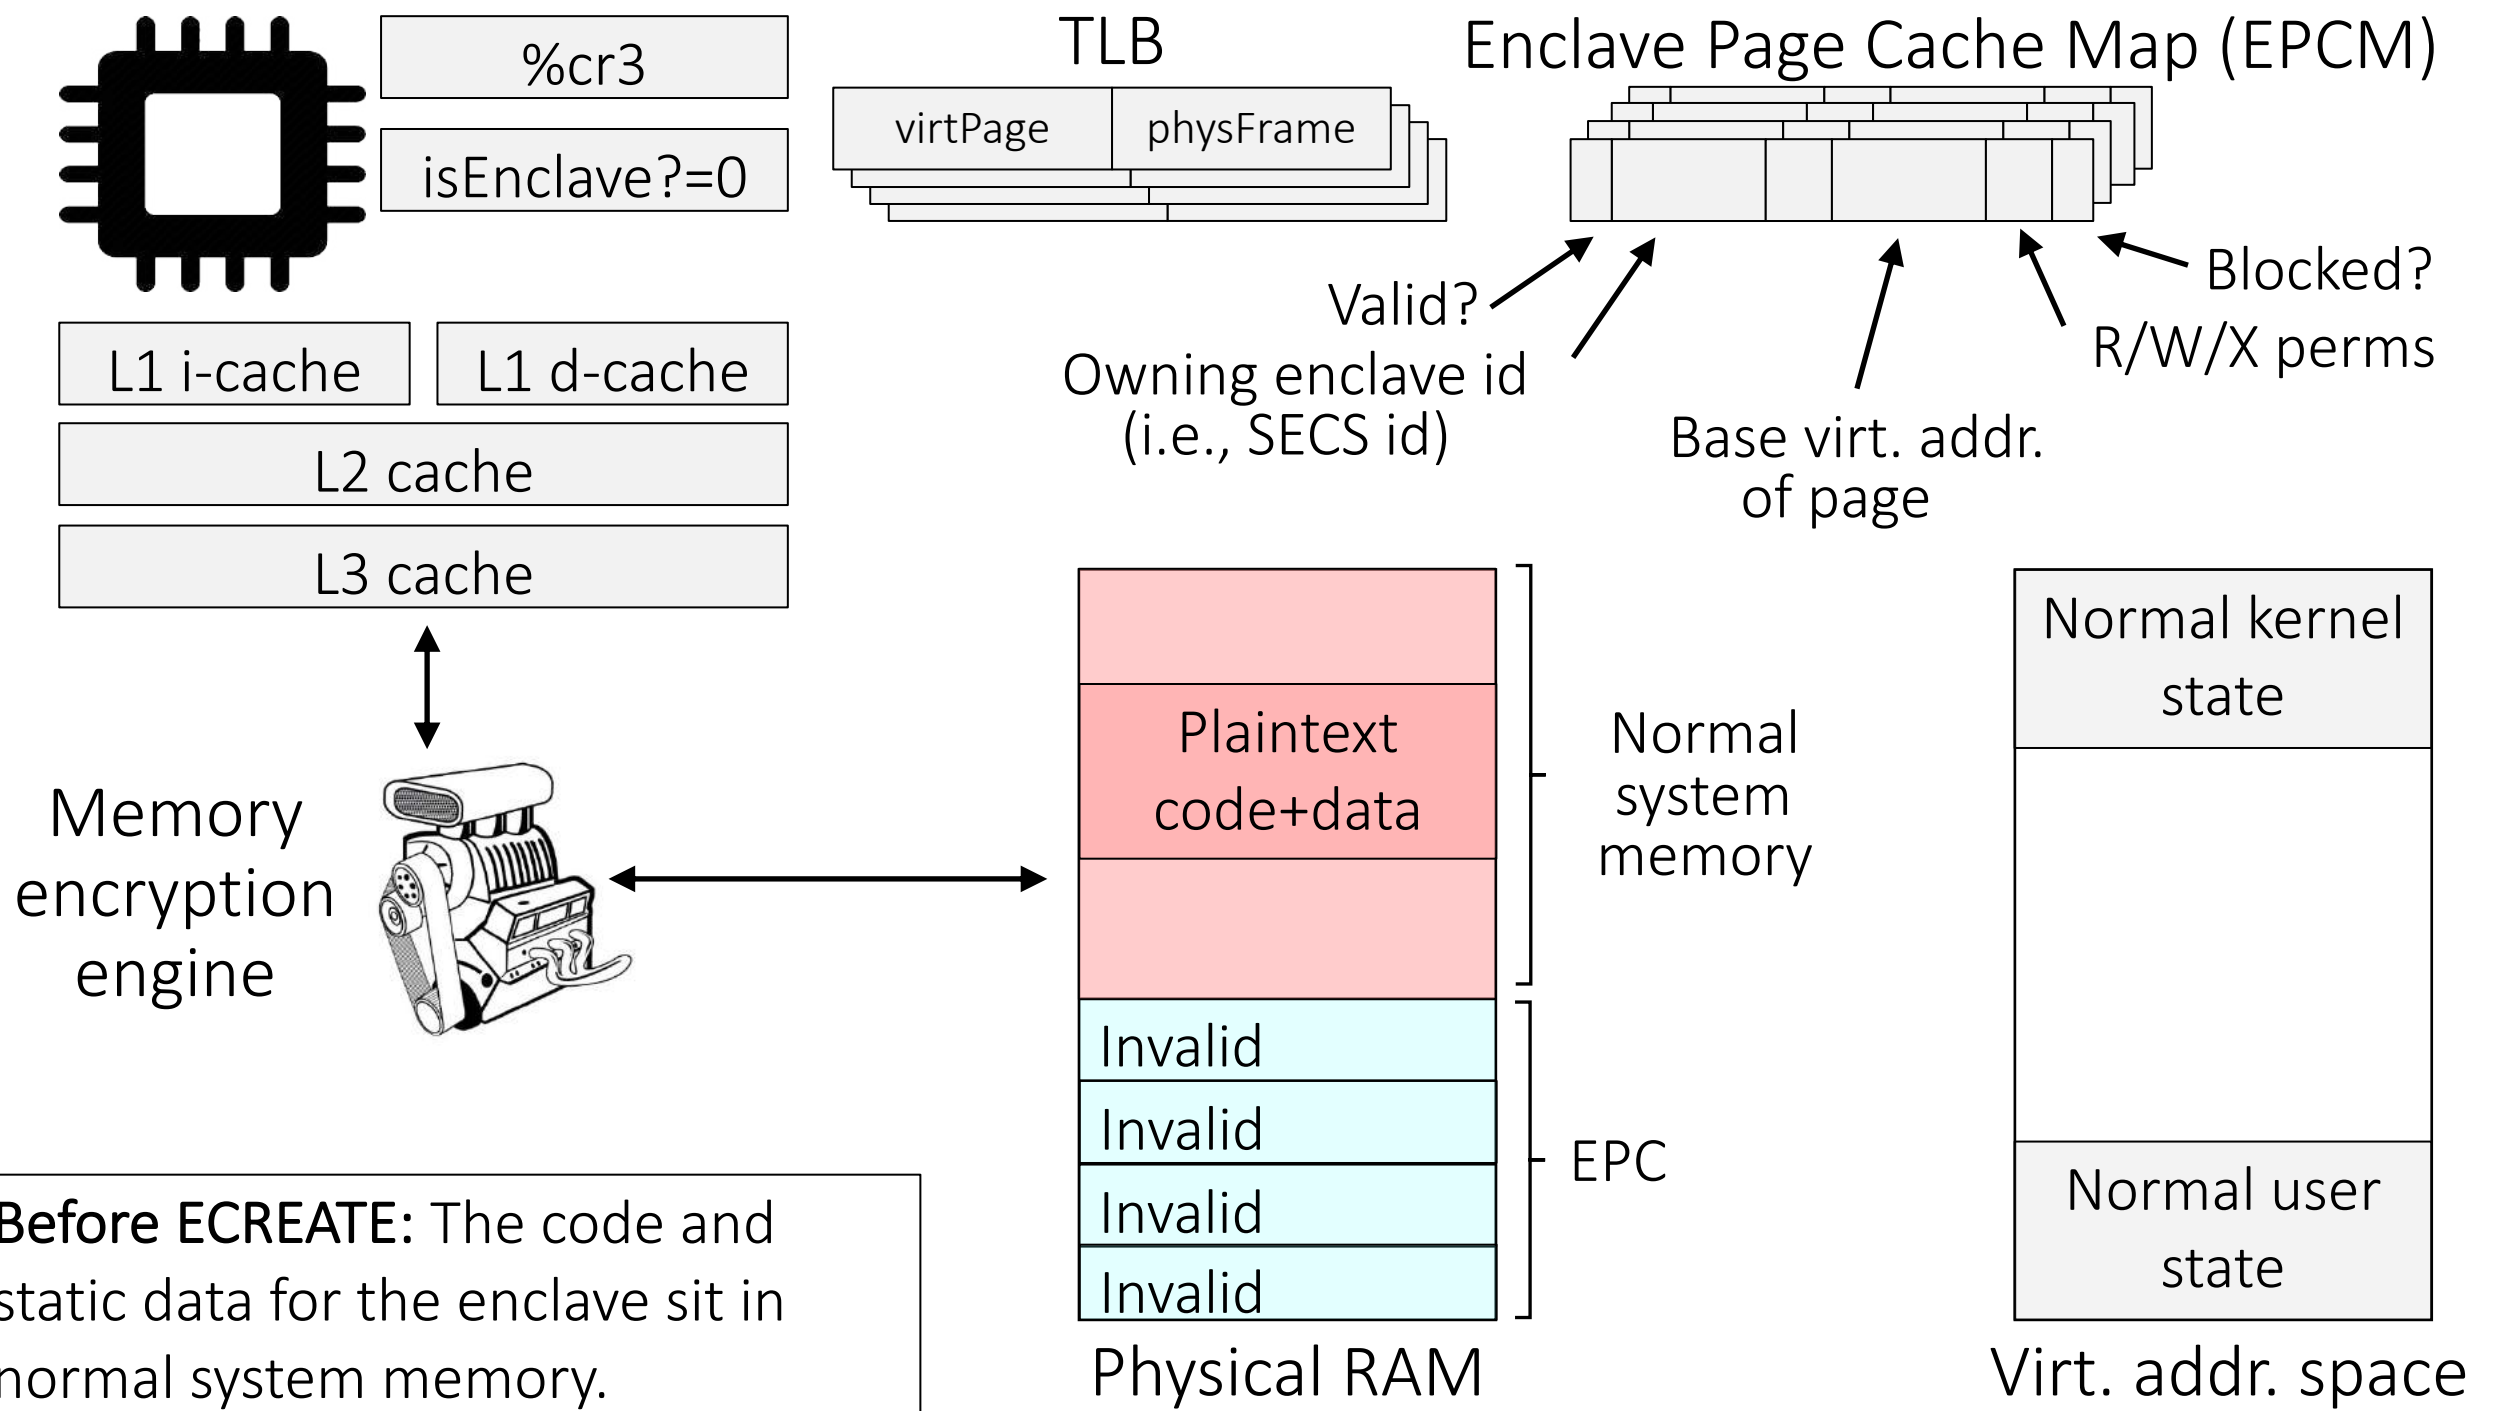
Normal
system
memory

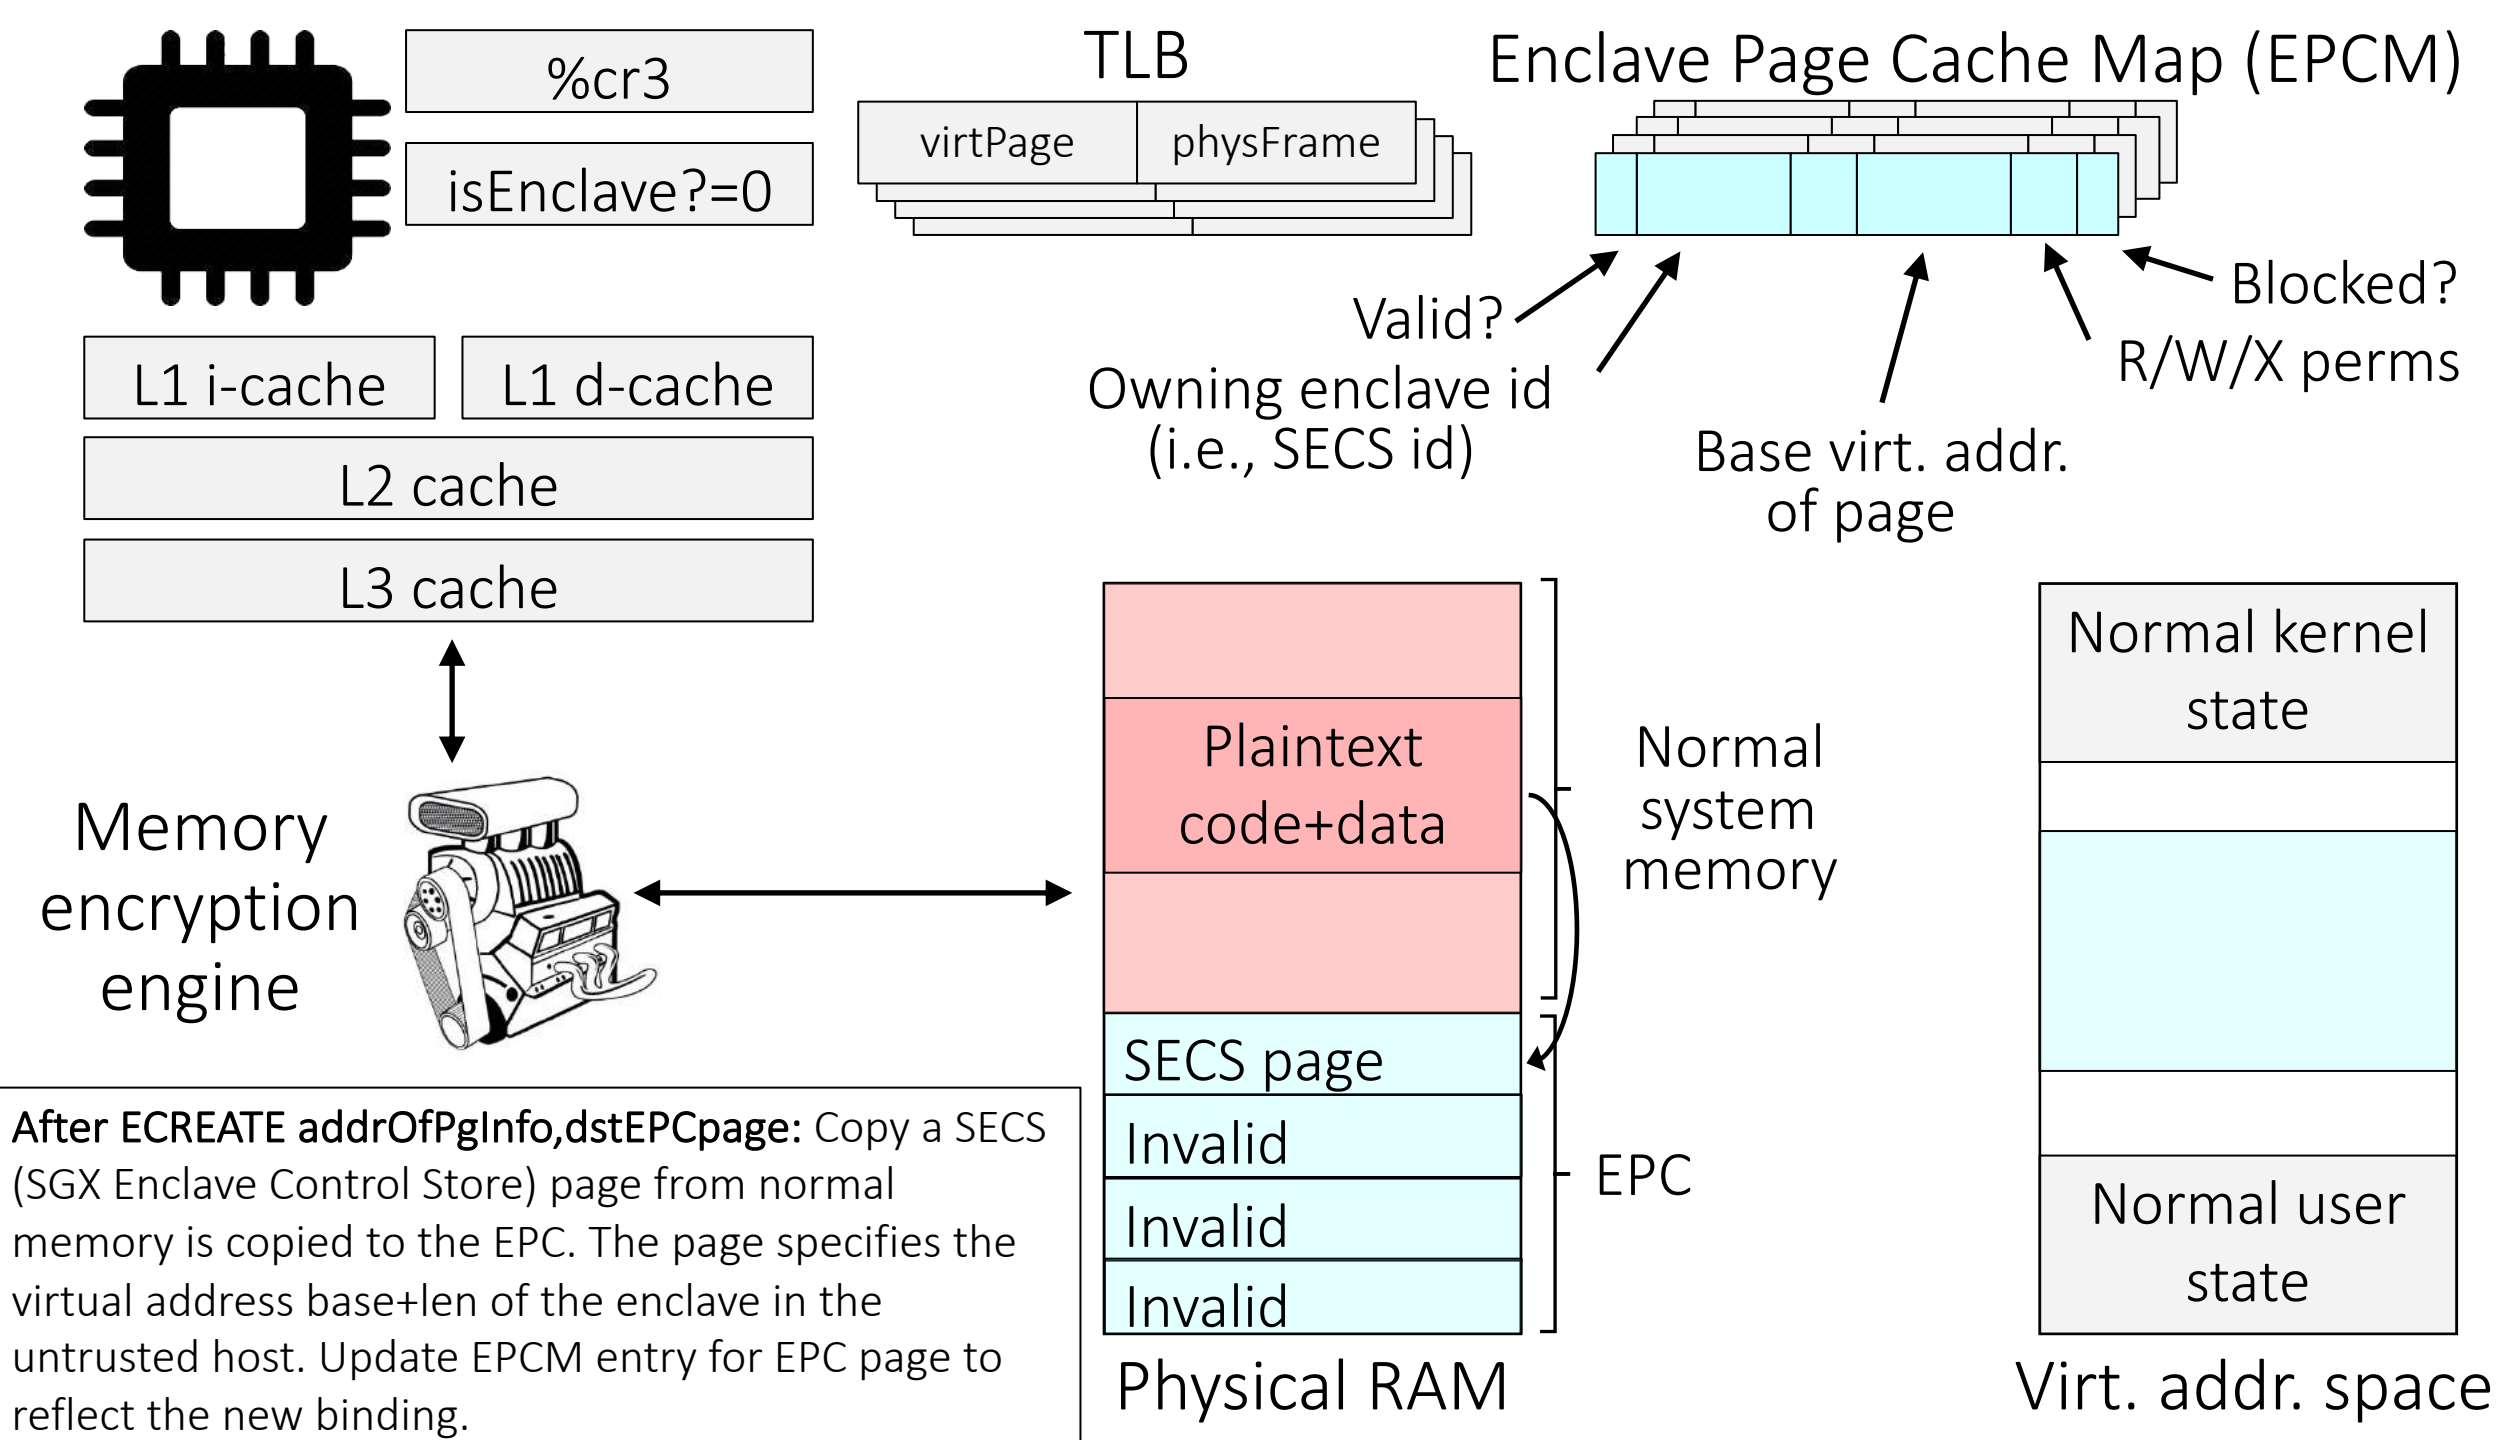
Enclave page
cache (EPC)

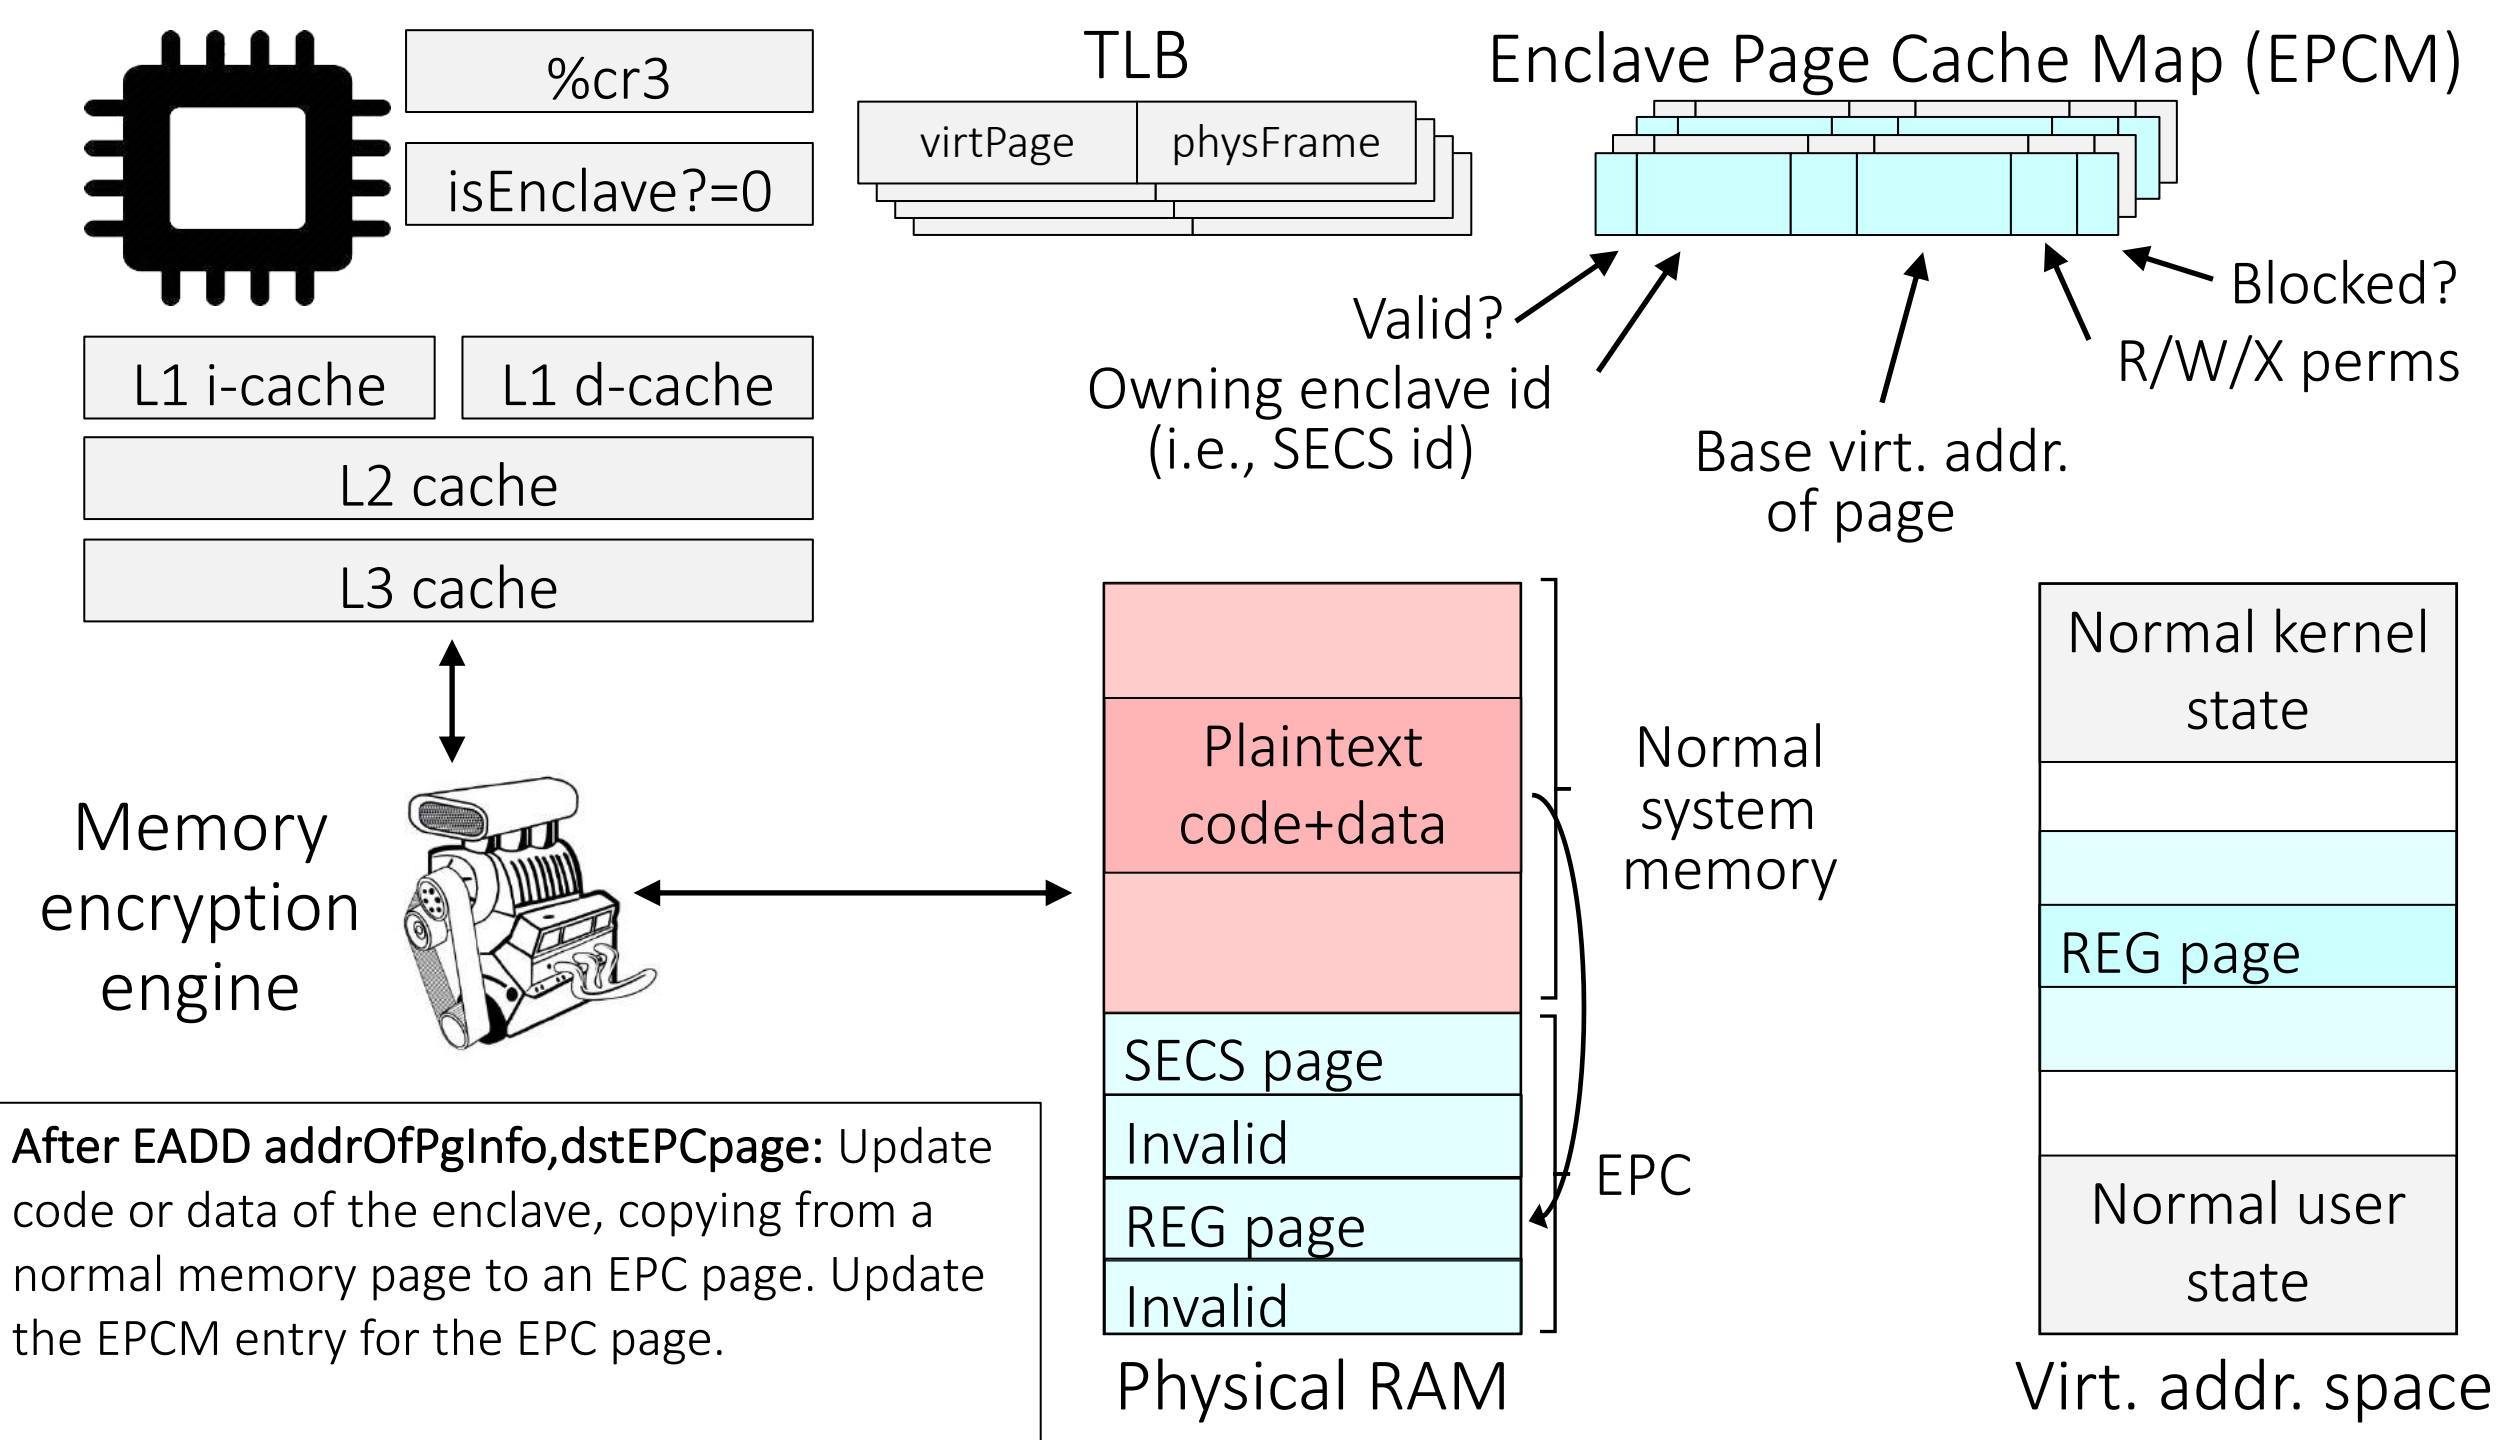
Physical RAM

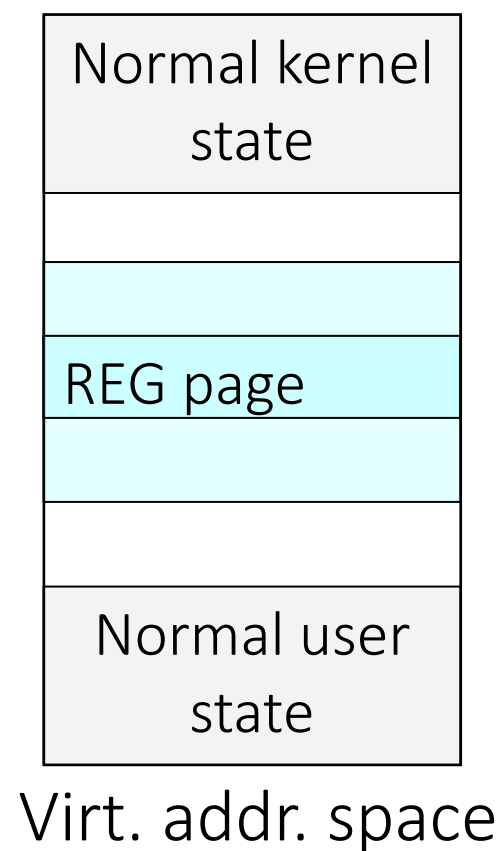
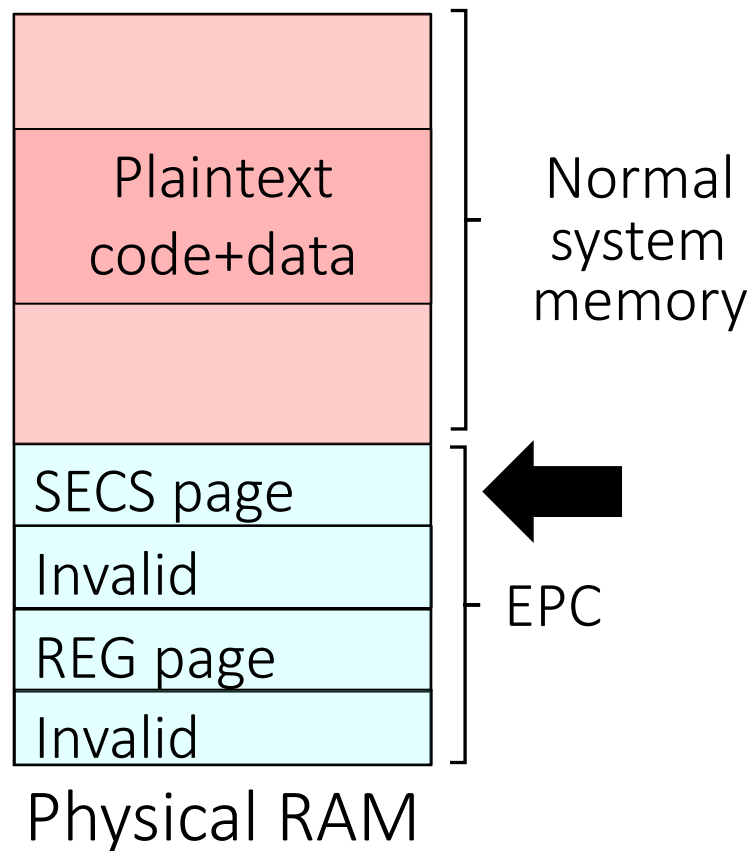
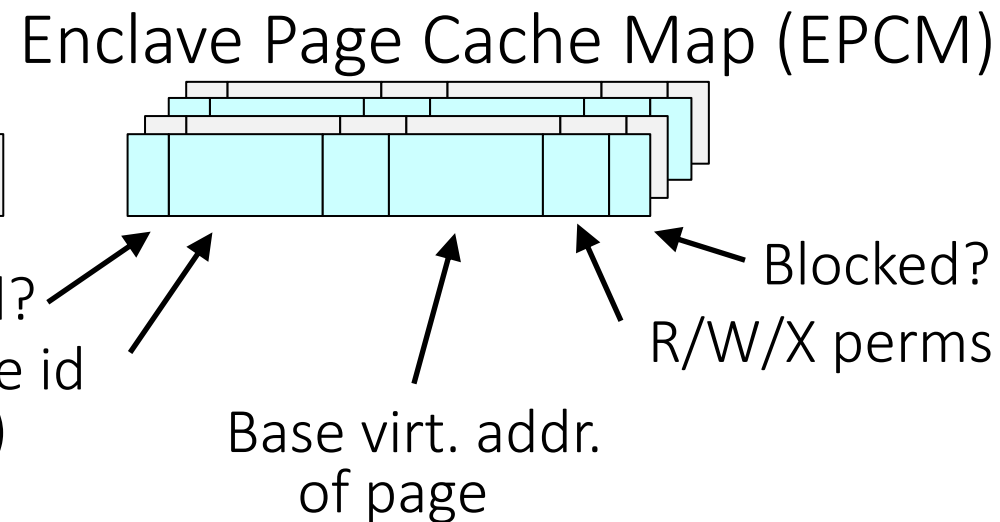




- EPCM is an SGX structure
 - Contains one metadata entry for each page in the EPC
 - Can only be modified by SGX instructions
 - Consulted during memory accesses to prevent EPC pages from unauthorized access

- Adds counters, MACs, and encryption to outgoing memory writes to EPC
- For incoming reads of EPC, decrypts data and then checks it for integrity and freshness (i.e., no rollbacks)










After EEXTEND addrOfSECSpage,
baseAddrOfEPCdataToHash:
Extend the hash value in the SECS page with 256
bytes of EPC data at the given address; do this
16 times to cover a4KB page.

SGX: EEXTEND

- By invoking **EEXTEND**, we build up a cumulative hash that can later be used to attest the state used to initialize an enclave
- Why does **EEXTEND** only cover 256 bytes instead of a full 4KB page?
 - Hashing an entire page would take longer!
 - A long-running instruction either has to:
 - Disable interrupts (which means that the system becomes less responsive to IO events, timers, etc.)
 - Allow interrupts but fail with an error code if interrupted (meaning that the instruction must be retried by the program that invoked it)
- Note that **ECREATE**, **EADD**, **EEXTEND**, and **EINIT** can only be called by privileged code (i.e., Ring 0 code)
 - OS can launch denial-of-service by refusing to load an enclave, or by loading the enclave improperly (e.g., omitting an **EEXTEND**)



%cr3

isEnclave?=0

L1 i-cache

L1 d-cache

L2 cache

L3 cache

Memory encryption engine



TLB

virtPage

physFrame

Enclave Page Cache Map (EPCM)

Valid?

Owning enclave id
(i.e., SECS id)

Base virt. addr.
of page

Blocked?

R/W/X perms

Plaintext
code+data

Normal
system
memory

SECS page

REG page

REG page

Invalid

EPC

Physical RAM

Normal kernel
state

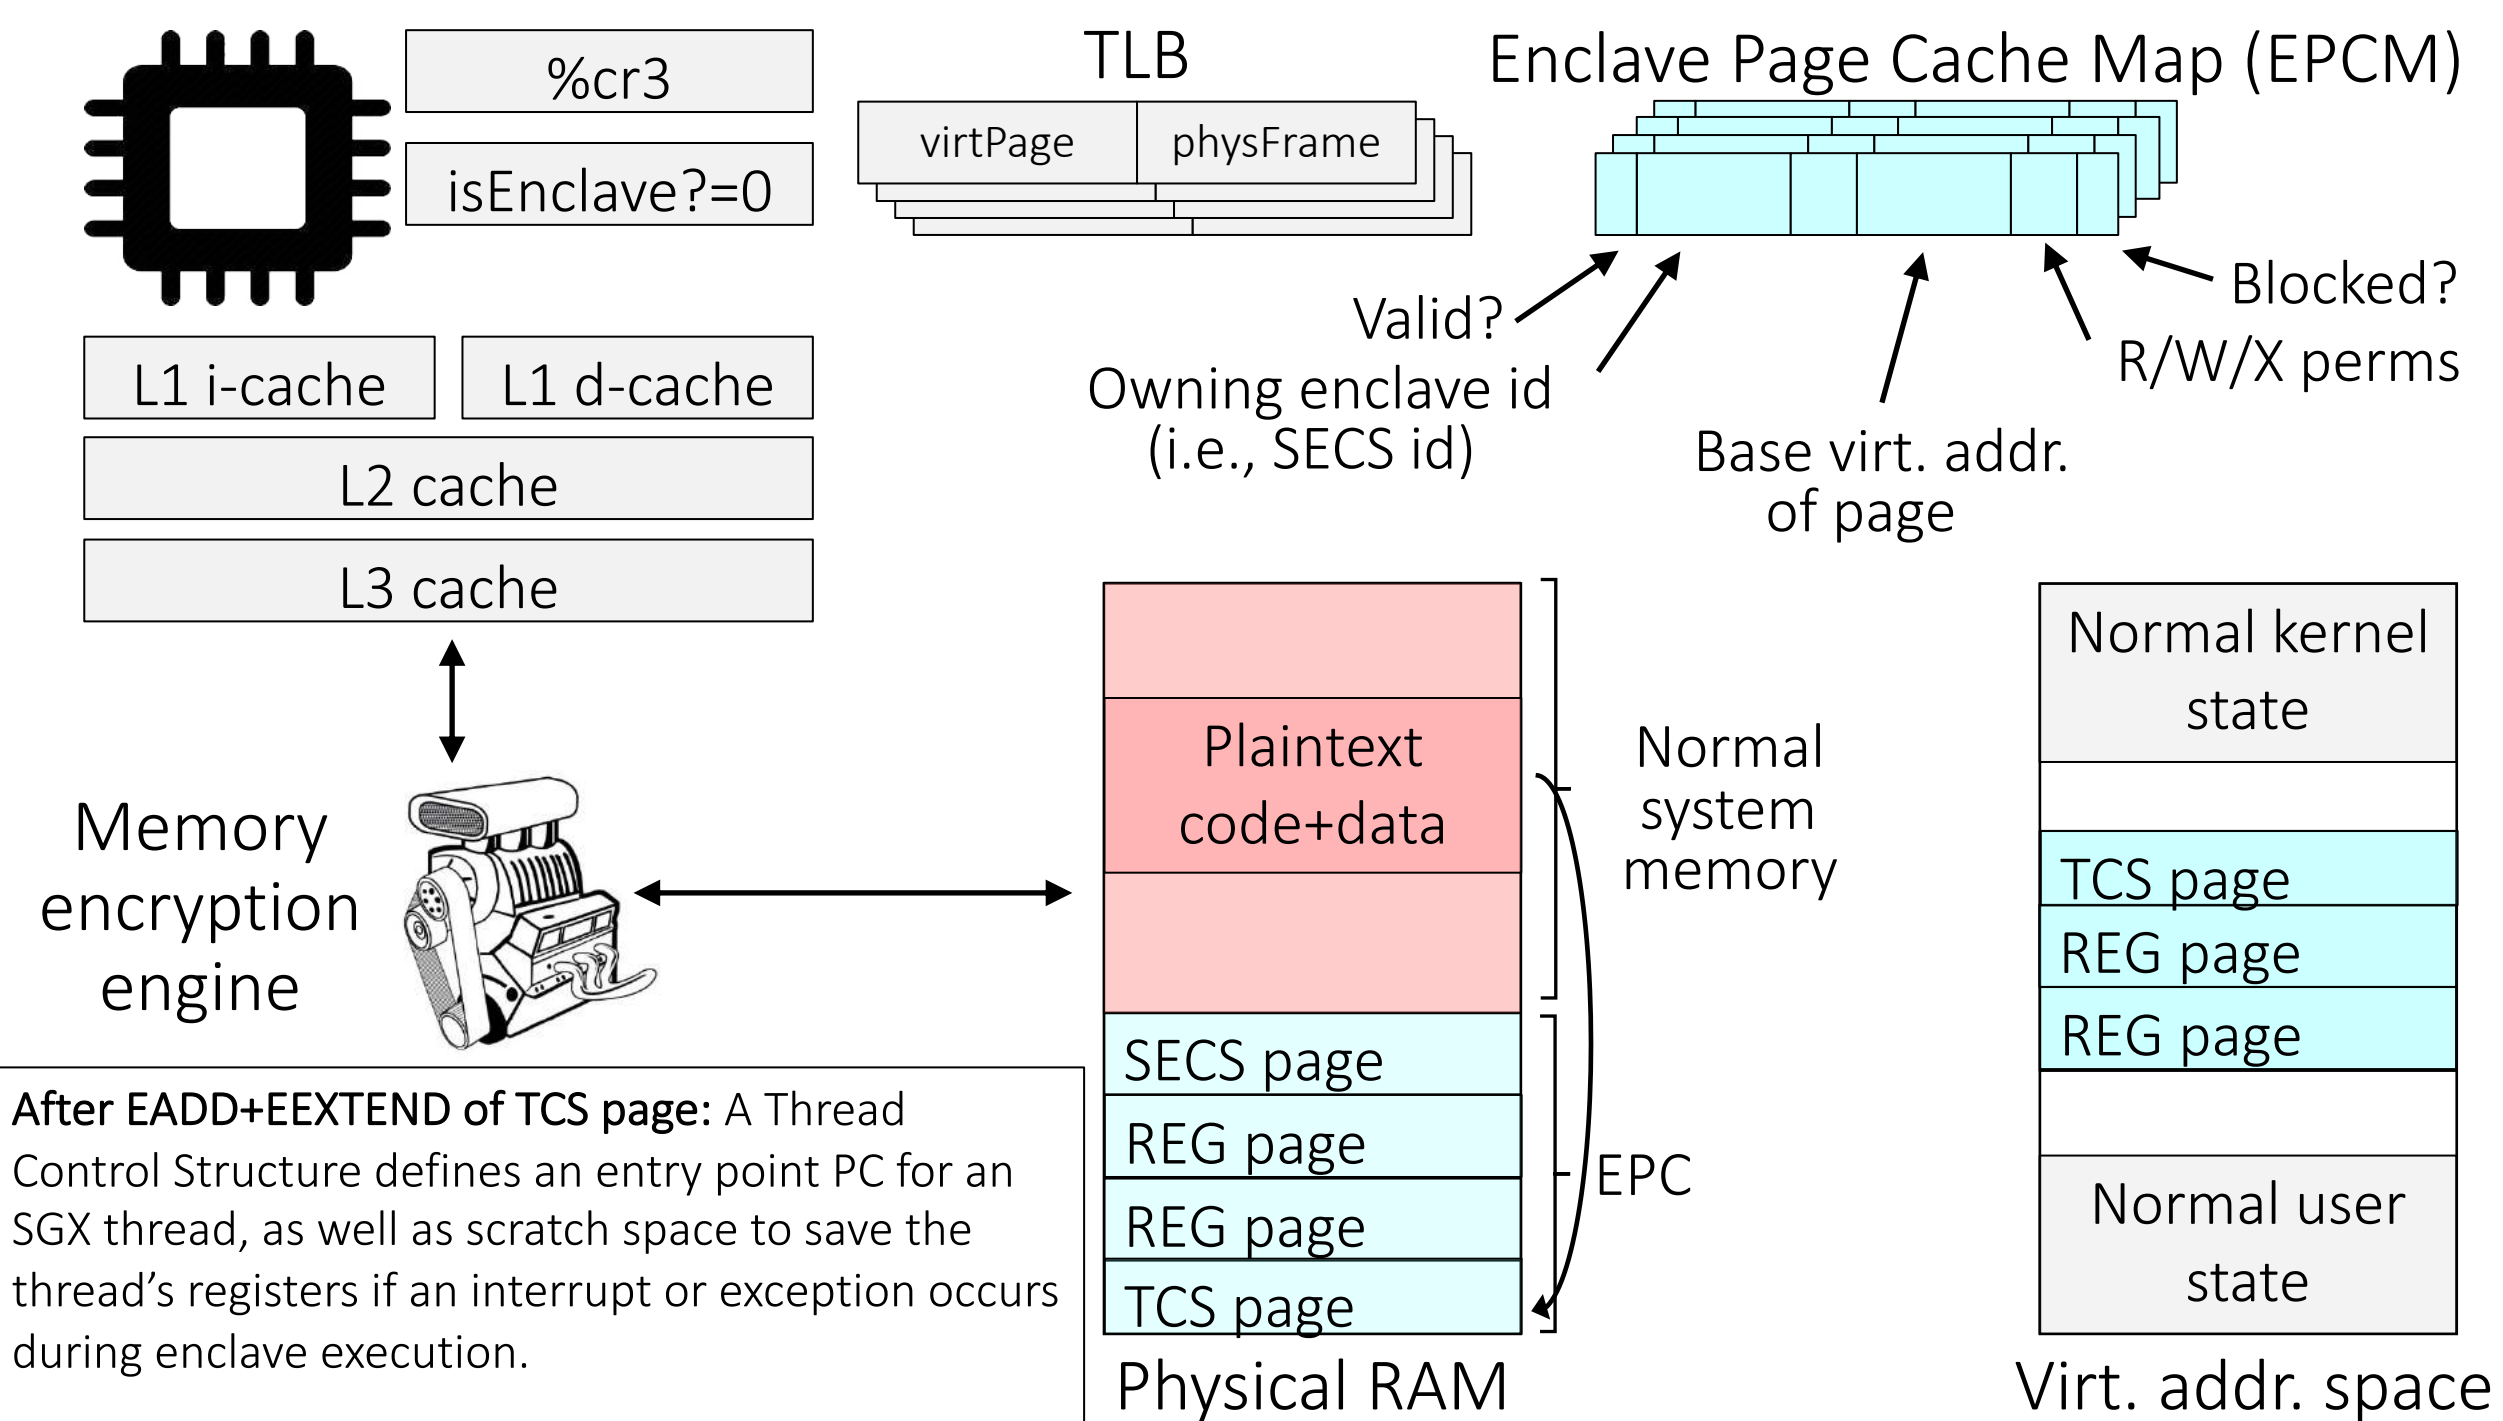
REG page

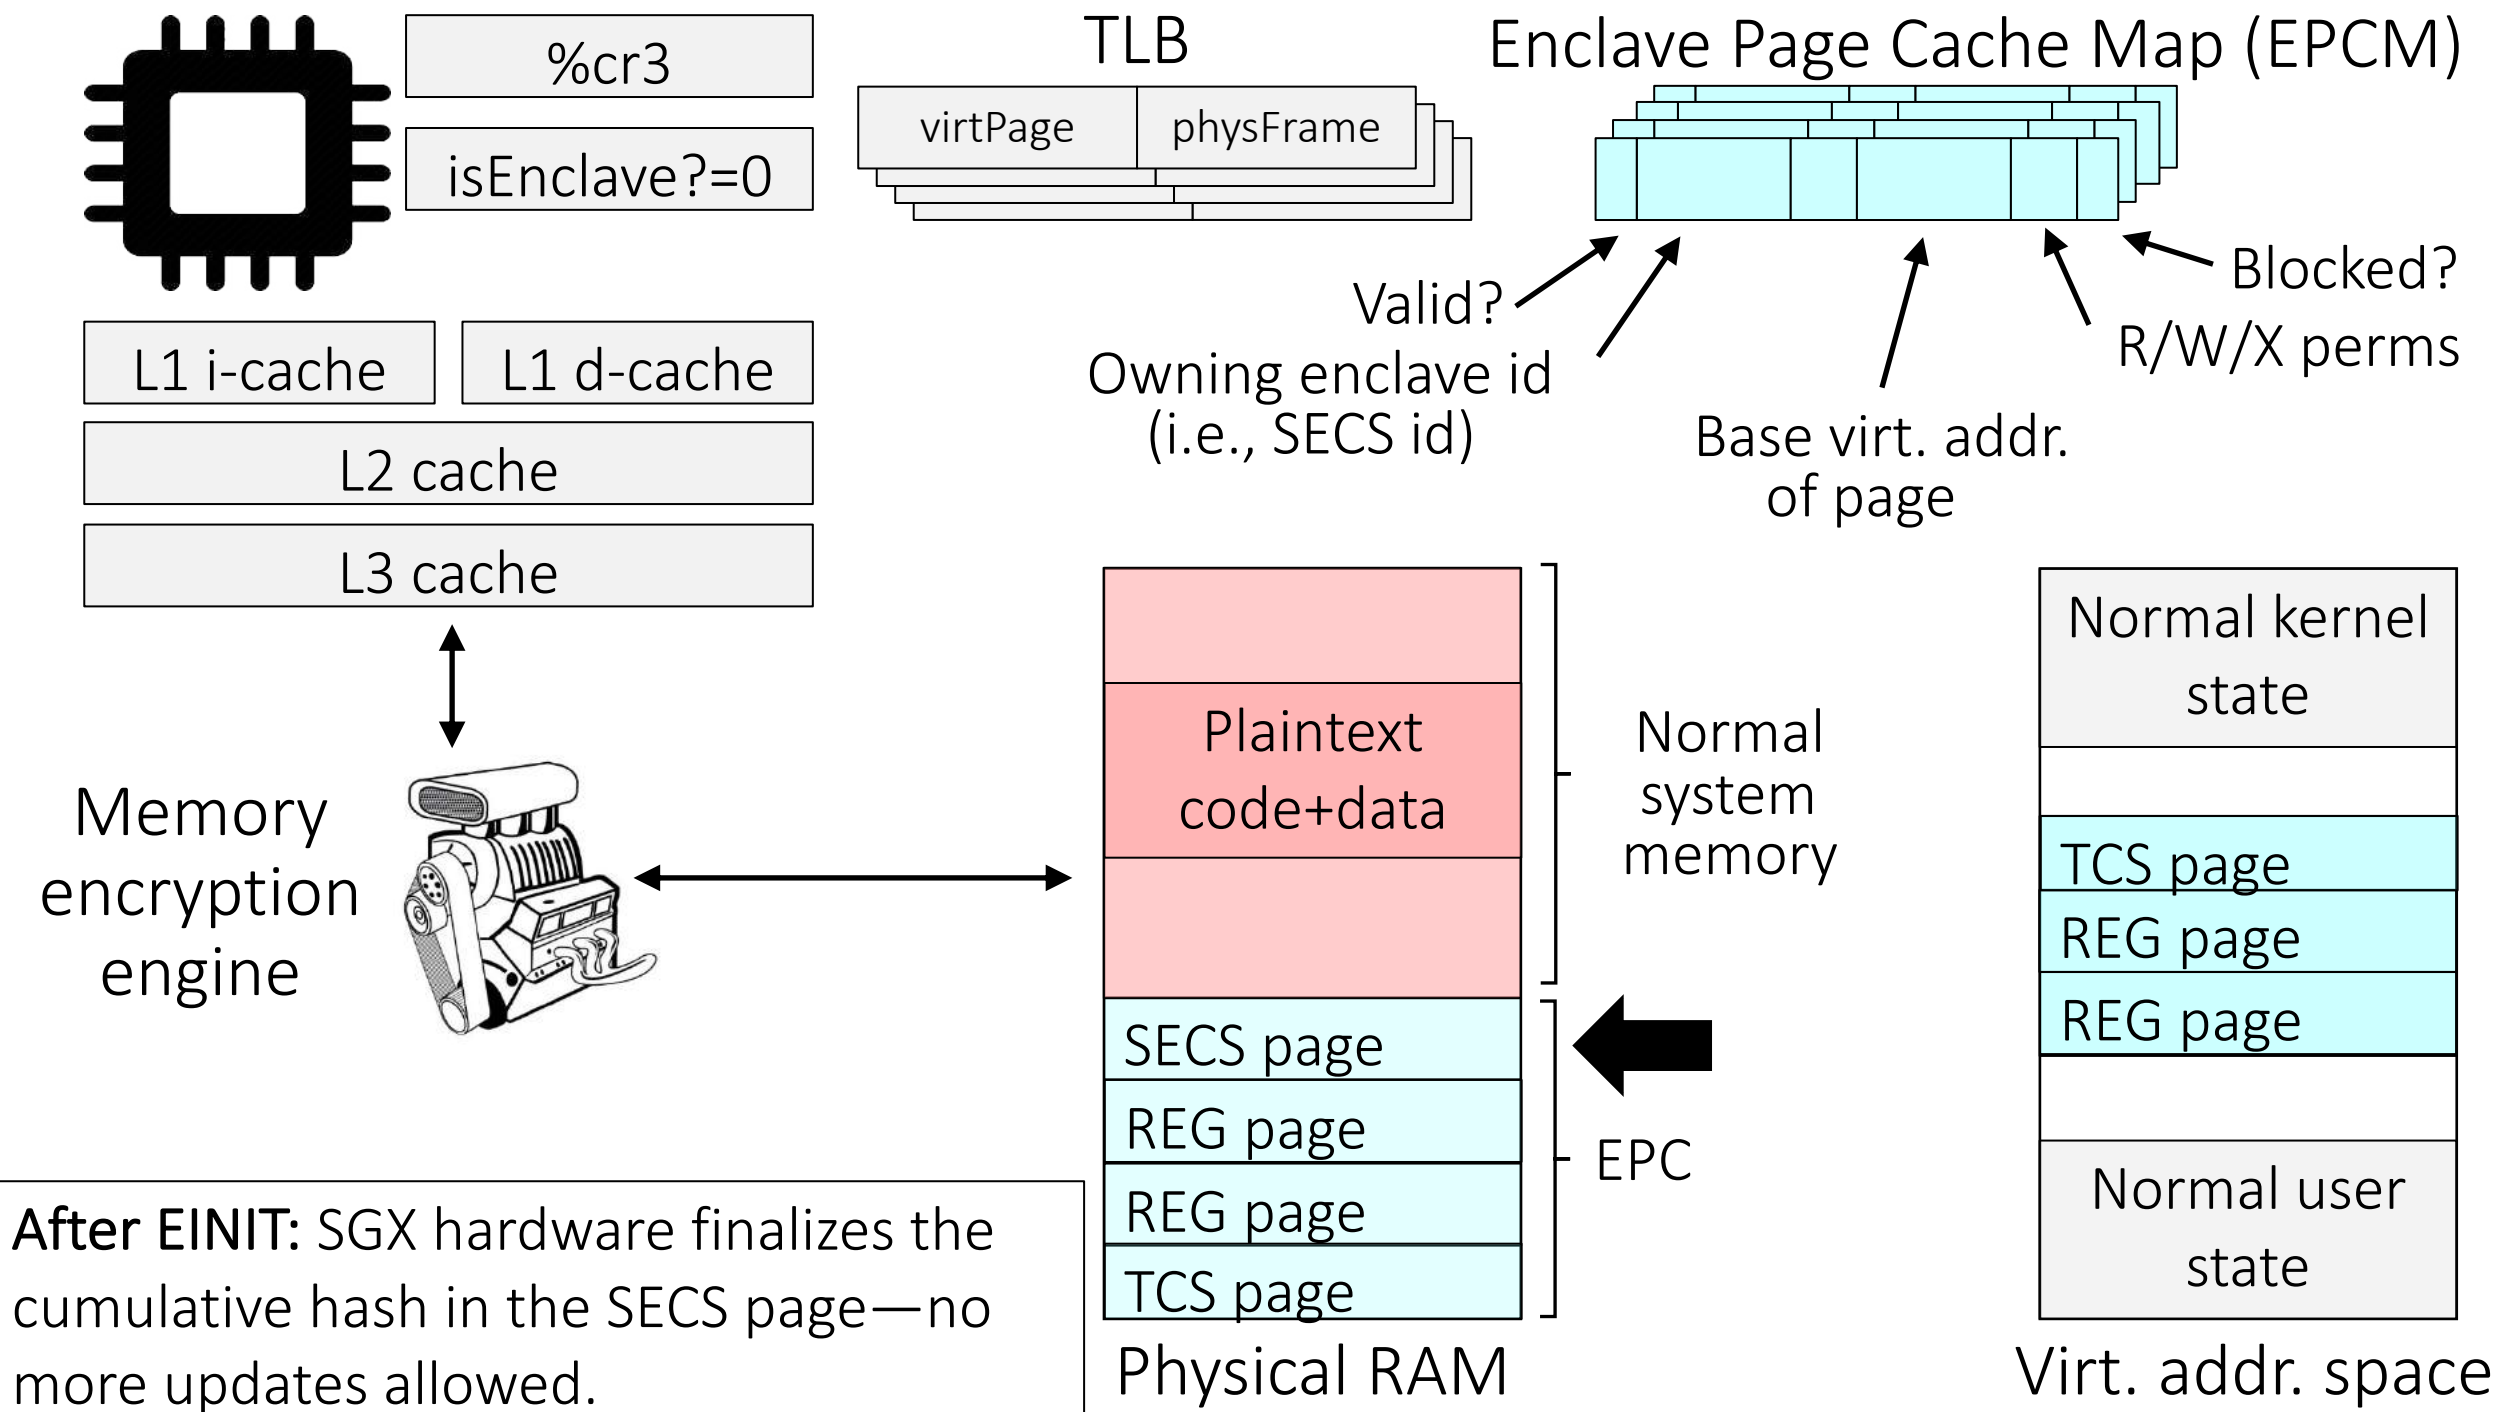
REG page

Normal user
state

Virt. addr. space

After another EADD and EEXTEND:
The EPCM is updated; additionally, the
cumulative hash in the SECS page is
updated.





EENTER and EEXIT

Neither instruction
changes %cr3!

- **EENTER** `addrOfTCS, AEP`

- Check **TCS.isBusy** to ensure that the untrusted host is not currently executing this SGX thread
- Flush TLB entries
- Flip **isEnclave** bit to 1
- Save **%rip**, **%rsp**, and **%rbp** to scratch area inside the enclave's address space
- Save **AEP** ("Asynchronous Exit Pointer" that points to a location in the untrusted host) to scratch area inside the enclave's address space
- Write the address of the instruction after **EENTER** to **%rcx**; this value should be saved by the enclave code
- Jump to the enclave start address described by TCS

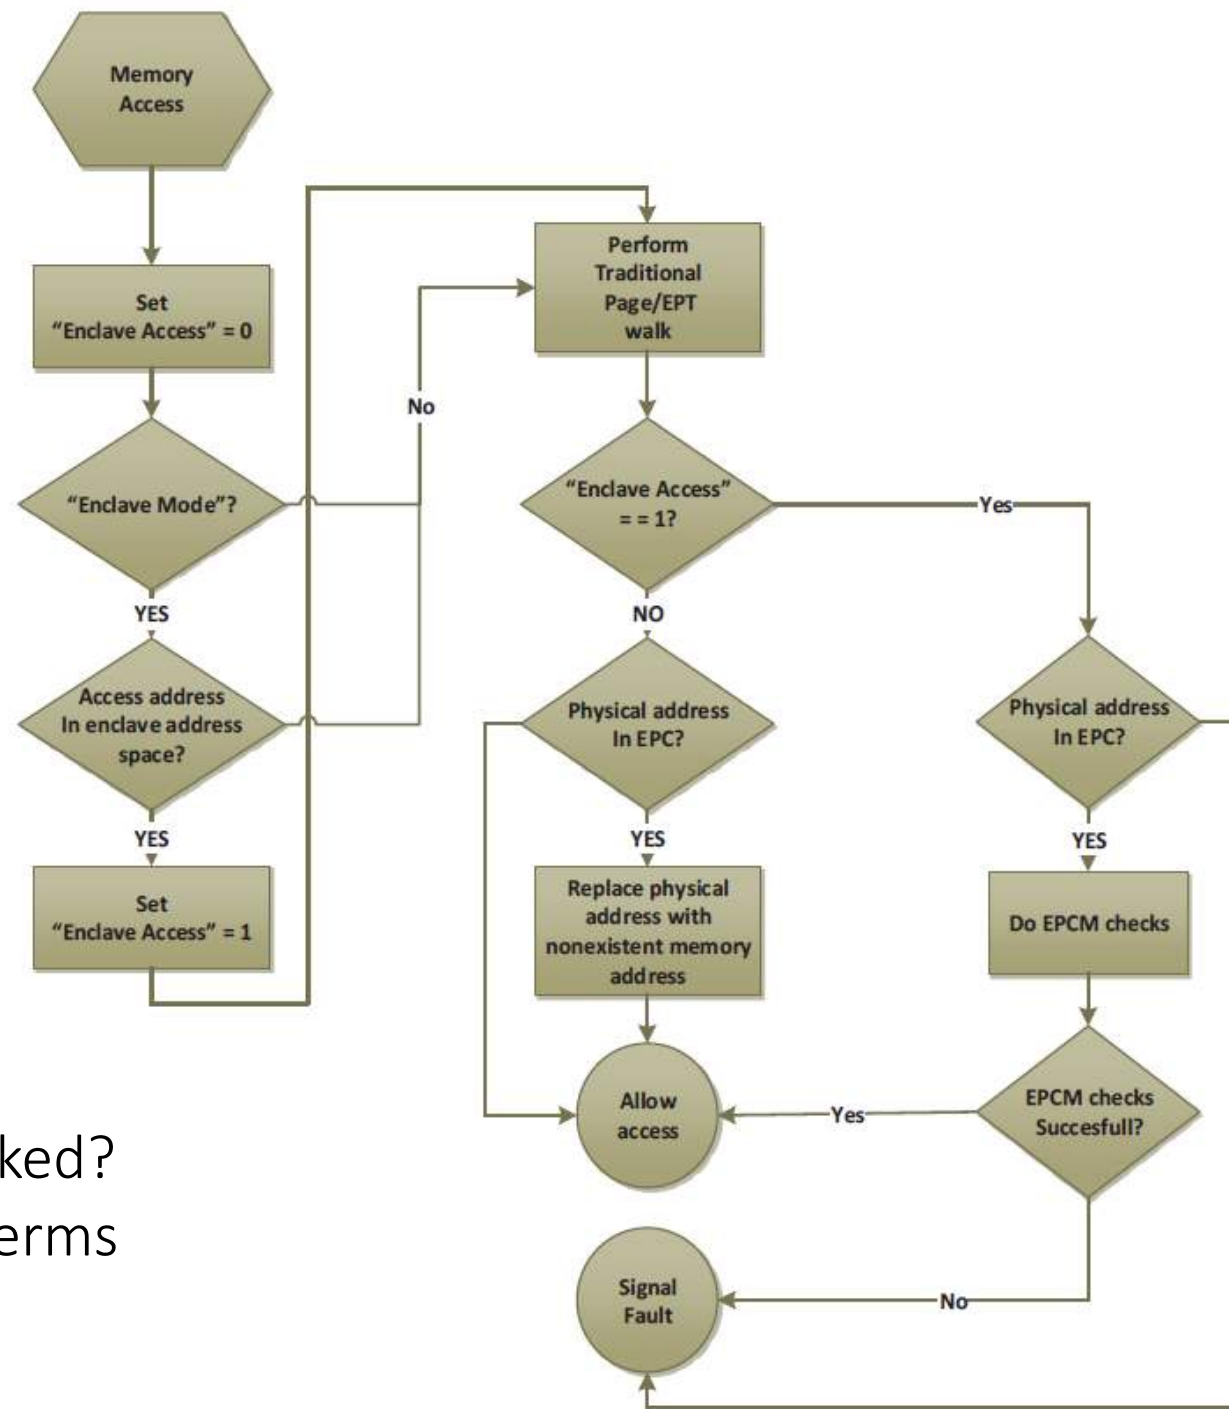

- **EEXIT** `jmpTarget`

- Flip **isEnclave** bit to 0
- Flush TLB entries
- Clear **TCS.isBusy**
- Jump to **jmpTarget** in the untrusted host; should be address saved by **EENTER**

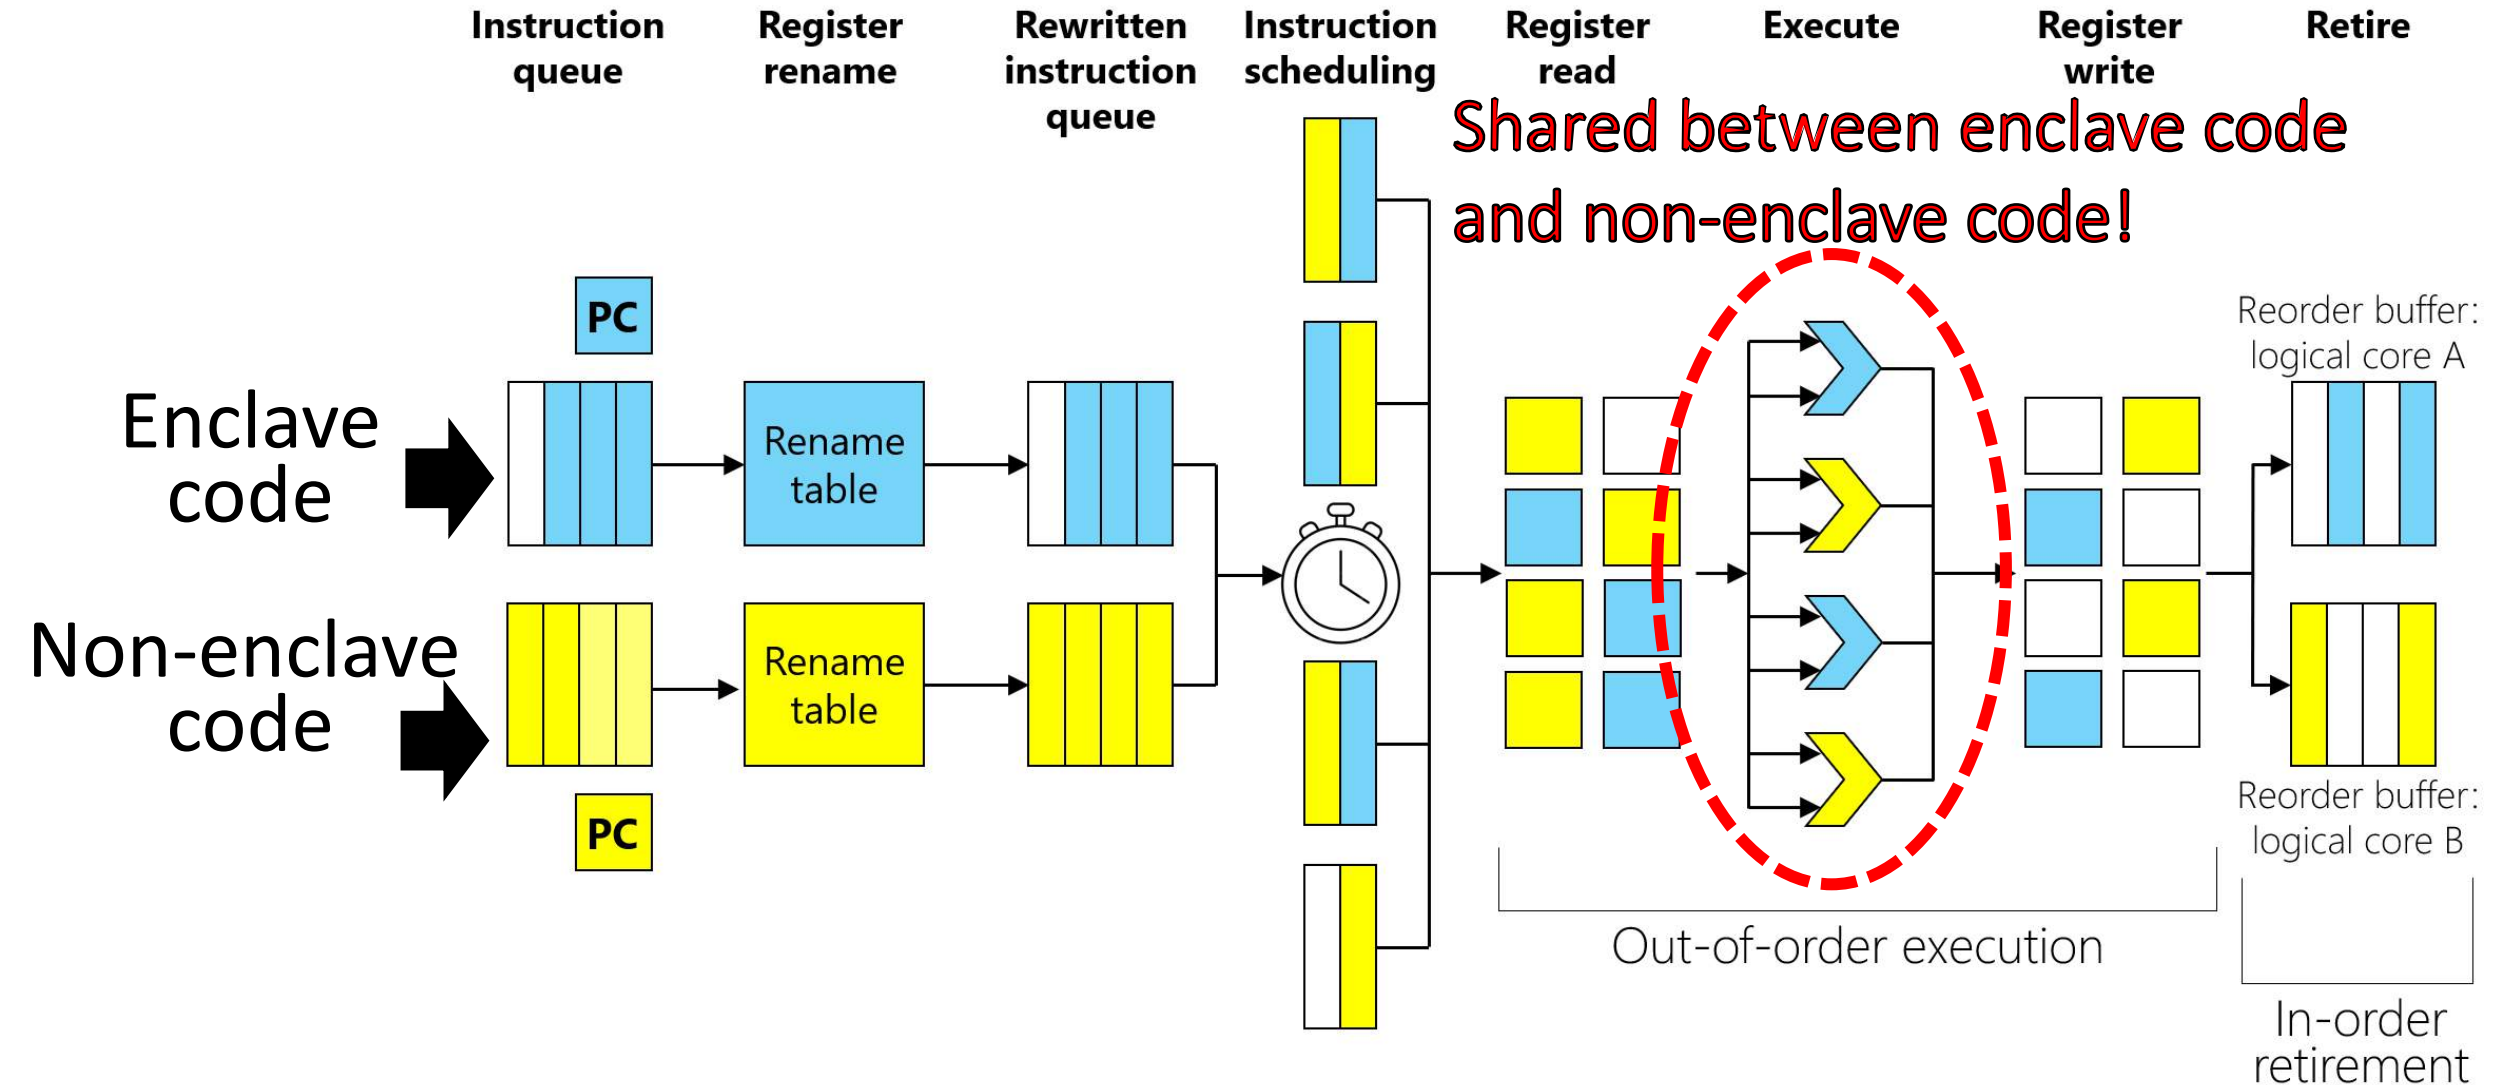
EENTER only callable by Ring 3+isEnclave=0
EEXIT only callable by Ring 3+isEnclave=1

Memory Access Checks in SGX

- SGX adds extra hardware to the memory controller to ensure that EPC pages can only be accessed by enclave code
- Extra hardware relies on EPCM



SGX is Vulnerable to Hyperthreading Side Channels!



SGX is Vulnerable to Cache Side Channels!

Rethinking Isolation Mechanisms for Datacenter Multitenancy

Varun Gandhi James Mickens
Harvard University

1 Introduction

Multitenancy is the foundation of modern cloud computing: a single datacenter machine must run code from multiple customers. To safely expose such a machine to untrusted tenants, datacenters have traditionally leveraged virtualization [4, 9]. In this approach, privileged hypervisor software (provided by the datacenter operator) mediates tenant access to raw physical resources like RAM and IO devices.

A hypervisor isolates tenants from each other, and isolates the hypervisor from tenants. However, tenants are not isolated from the datacenter operator; the operator's hypervisor can arbitrarily manipulate tenant state, and the operator herself can physically inspect or modify the contents of server RAM. Intel's SGX-enabled processors [13] stop these attacks. Using hardware-enforced memory partitioning, SGX prevents a hypervisor from accessing secure tenant pages. SGX hardware also transparently encrypts and HMACs cache lines during eviction to RAM; thus, a datacenter operator with physical control of a machine cannot see cleartext tenant RAM, or undetectably tamper with the encrypted RAM that *is* visible. SGX is the foundation for a variety of software-level runtimes that isolate datacenter tenants from privileged management software [3, 5, 45, 58].

Unfortunately, SGX-based approaches have three important limitations.

- SGX gives a tenant the illusion of ISA-level isolation. However, tenants cohabitate at the *microarchitectural* level, resulting in side channel vulnerabilities that leak information from ostensibly secure computations (§2.1).
- SGX can cryptographically vouch for the initialization-time integrity of a secure computation. However, SGX has no way to attest a computation's *dynamic* (i.e., *post-load*) integrity. Both initial and post-load integrity are important (§3.2). Clients do not wish to exchange

Motivated by these problems, we propose a new isolation approach for datacenter multitenancy. As with SGX, we leverage trusted hardware to isolate tenants from each other and from the datacenter operator. However, our approach differs from SGX in three crucial ways.

- First, our trusted hardware strongly isolates each tenant's ISA-level state at the microarchitectural level, removing side channels involving other tenants or the hypervisor.
- Second, we allow a tenant to explicitly bind application code to *monitor code* that dynamically enforces runtime security invariants like control flow integrity. The monitor code runs in parallel with application code; however, the monitor runs on a different CPU pipeline (managed by trusted hardware) that receives a read-only stream of the register state from the application-level pipeline. With the exception of this register mirroring, the two pipelines are isolated at the microarchitectural level. This design prevents side channel leakages of monitor state to application code that might be under attack. Microarchitectural partitioning of an application and its monitor also eliminates more direct attacks that could occur if monitor state were located in the same address space as the application to protect.
- Third, our new trusted hardware uses an open microcode format, and exposes a software-readable description of microarchitectural-level hardware details. This approach allows tenants to independently verify the security properties of a server's hardware. This design also allows tenants to customize monitor code to fully exploit the microarchitectural affordances provided by a particular datacenter server.

In Section 2, we provide more background on SGX and related technologies. We then sketch a preliminary design for our new isolation hardware (§3). We conclude by describing

- The enclave and the untrusted host share the same L1 cache!
- Untrusted host can:
 - Load the enclave and evict an oracle array's cache lines
 - Wait for the enclave to bring a cleartext secret value into the L1 cache, then do:

```
uint8_t v = *secret_ptr;
uint64_t o = oracle[v*4096];
```
 - The read will fail at the ISA level, but at the microarchitectural level, a speculative load will read a cache line from the oracle array!
 - The untrusted host can then time how long it takes to read each cache line in the oracle array, determining the secret byte!

What Did We Learn?

Day I: Introduction to Control Flow Integrity

- Overview of virtual address spaces
- Attacks: Buffer overflows, return-to-libc
- Defenses: NX bits, Intel shadow stacks, ASLR

Day II: Advanced Attacks and Defenses

- Attacks: Memory vulnerabilities via type confusion, ROP attacks, JOP attacks
- Defenses: Compiler-enforced CFI, Intel indirect branch tracing
- SGX: Design and vulnerabilities

Software-level Attacks on Architectural and Microarchitectural State



ACACES SUMMER SCHOOL

July 9th - 10th, 2020

~~Fiuggi, Rome~~ Zoom World

James Mickens
Harvard University